

# Quels facteurs de rapidité de calcul ?

- Computing unit : nombre d'opérations par seconde
- Memory unit : vitesse d'accès mémoire
- Communication : vitesse de transfert des données

# Coté processeur (computing unit)

- Ce qu'on veut améliorer:
  - Le nombre d'opérations par secondes
- Deux principaux paramètres de contrôle :
  - nombre d'opérations (instructions) par cycle (IPC)
  - nombre de cycles par seconde (clock speed)

- Limitations :
  - physiques, transistors trop petits, chauffage trop important
  - Consommation électrique : fréquence<sup>3</sup>

	mono-cœur	bi-cœurs
Fréquence	F	0.75F
Consommation par processeur	W	0.84 W
Performance par processeur	P	1.5P

# Coté mémoire (memory units)

- Ce qu'on veut améliorer :
  - Diminuer la *latency* (temps d'accès à l'information, exprimée en temps ou nombre de cycles)

- Exemples :
  - Disque dur à tête de lecture
  - Solid state hard drive
  - RAM
  - L1/L2 cache
  - Registres ( $\sim 1$  cycle)

## Characteristics of various memory units

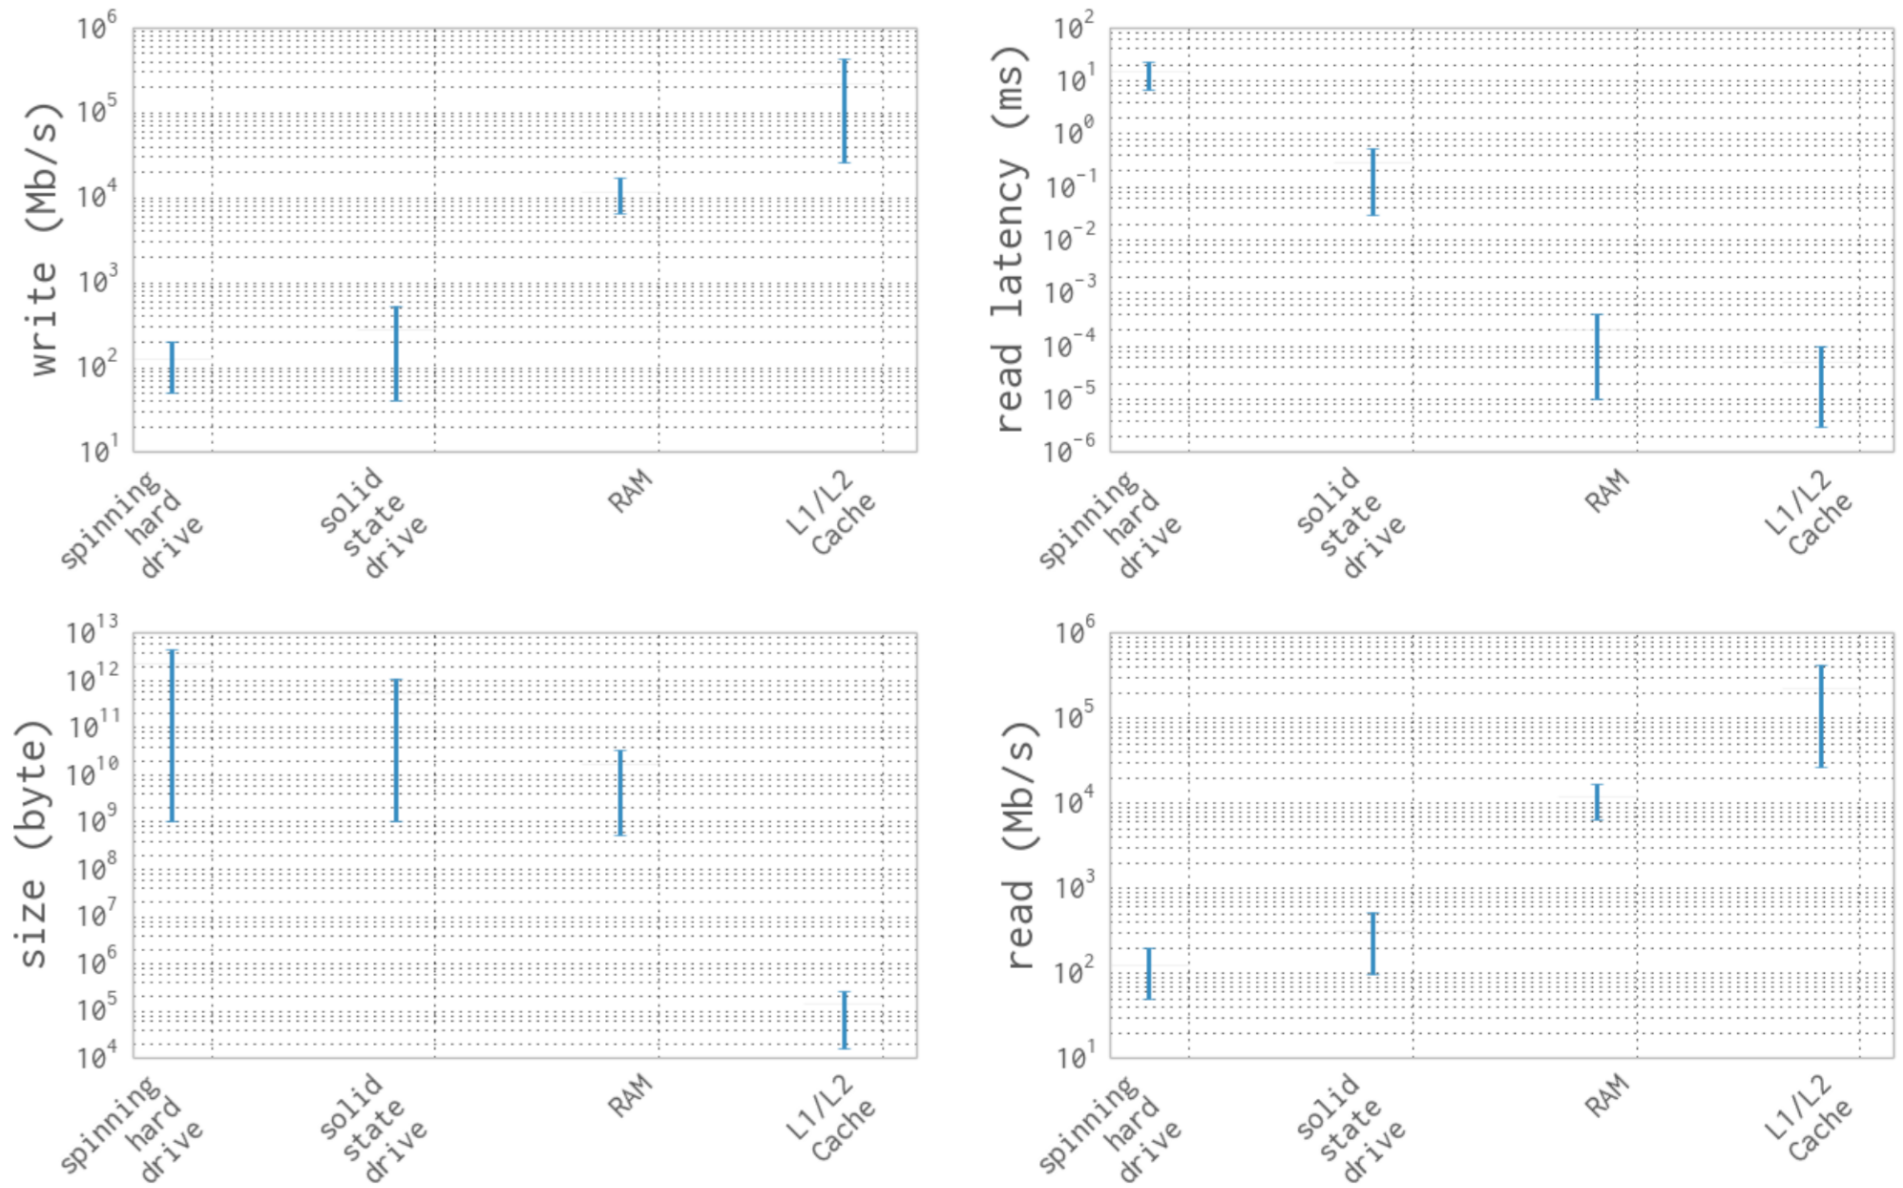
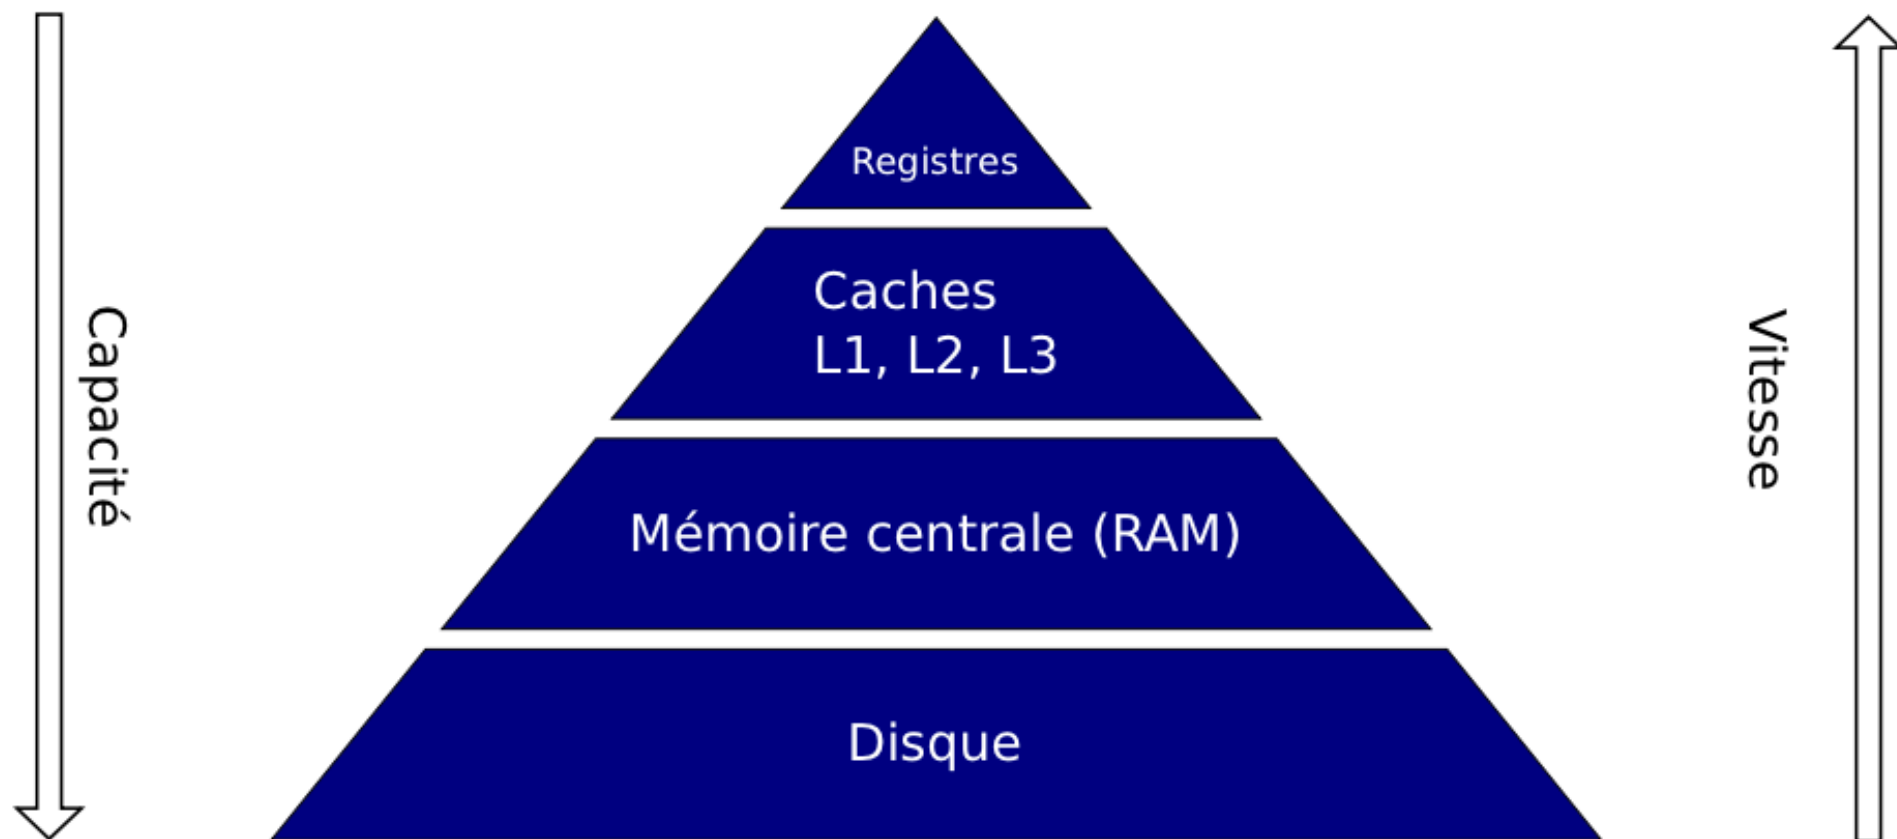


Figure 1-2. Characteristic values for different types of memory units (values from February 2014)



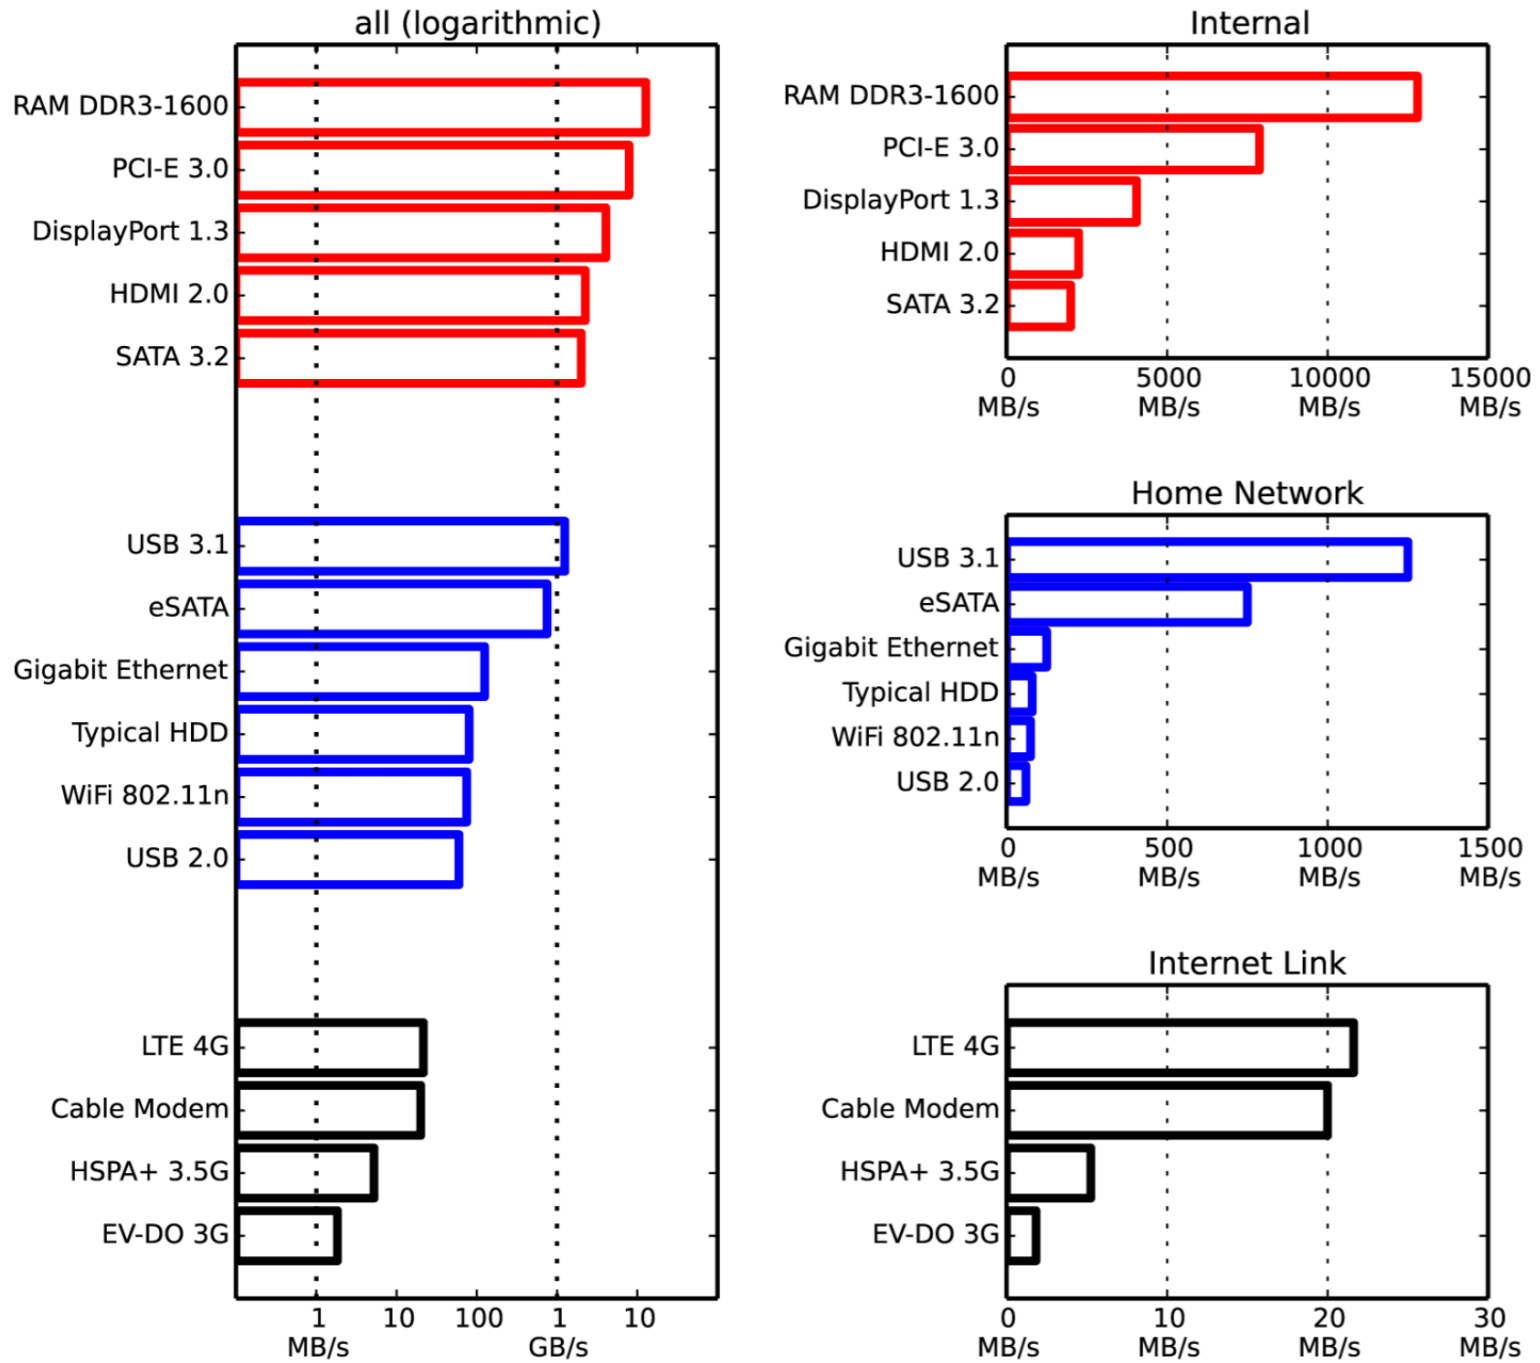
# Communication

- Ce qu'on veut améliorer :
  - Data transfer rate (*bits*) from memory components to computing units
- Two main parameters :
  - *Bus width* : amount of bit that can be moved at once
  - *Bus frequency* : how many times per second



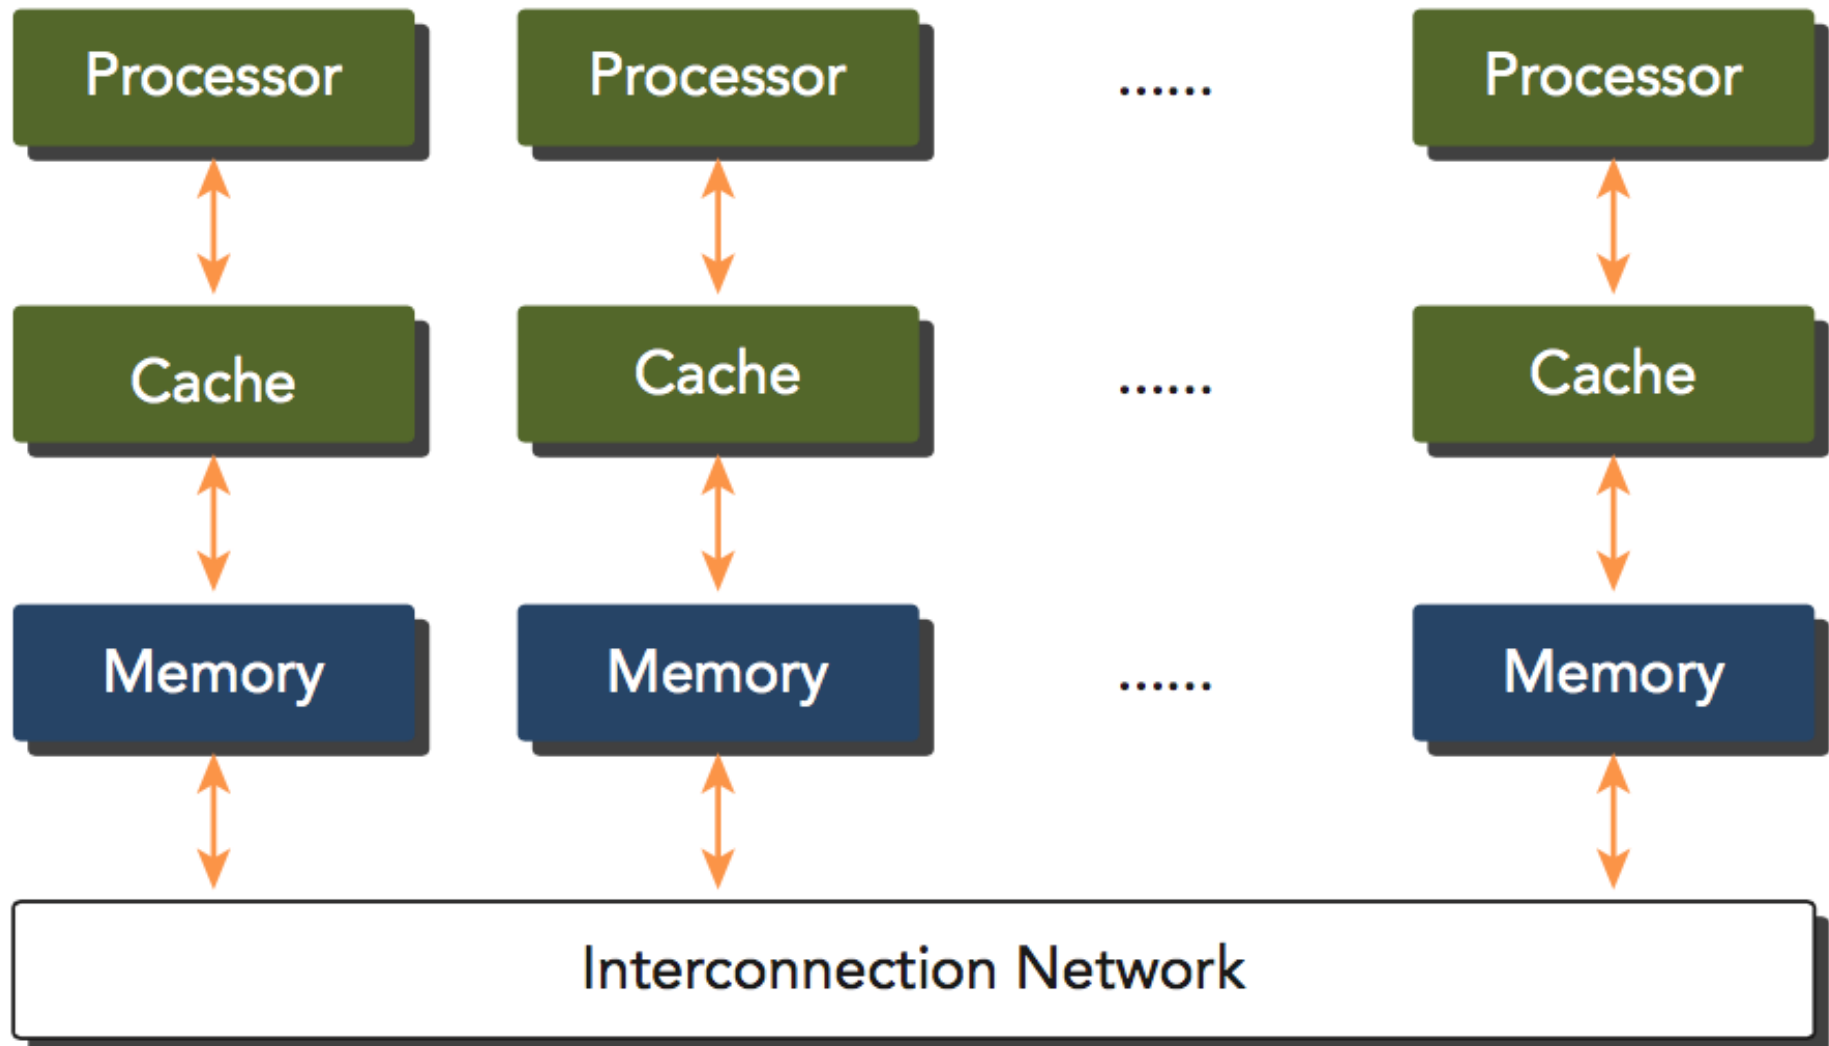
- Exemples :
  - frontside bus (between RAM et caches)
  - backside bus (between caches et CPU)
  - external bus (autres périphériques, carte réseau...)

# Bandwidth of Common Interfaces



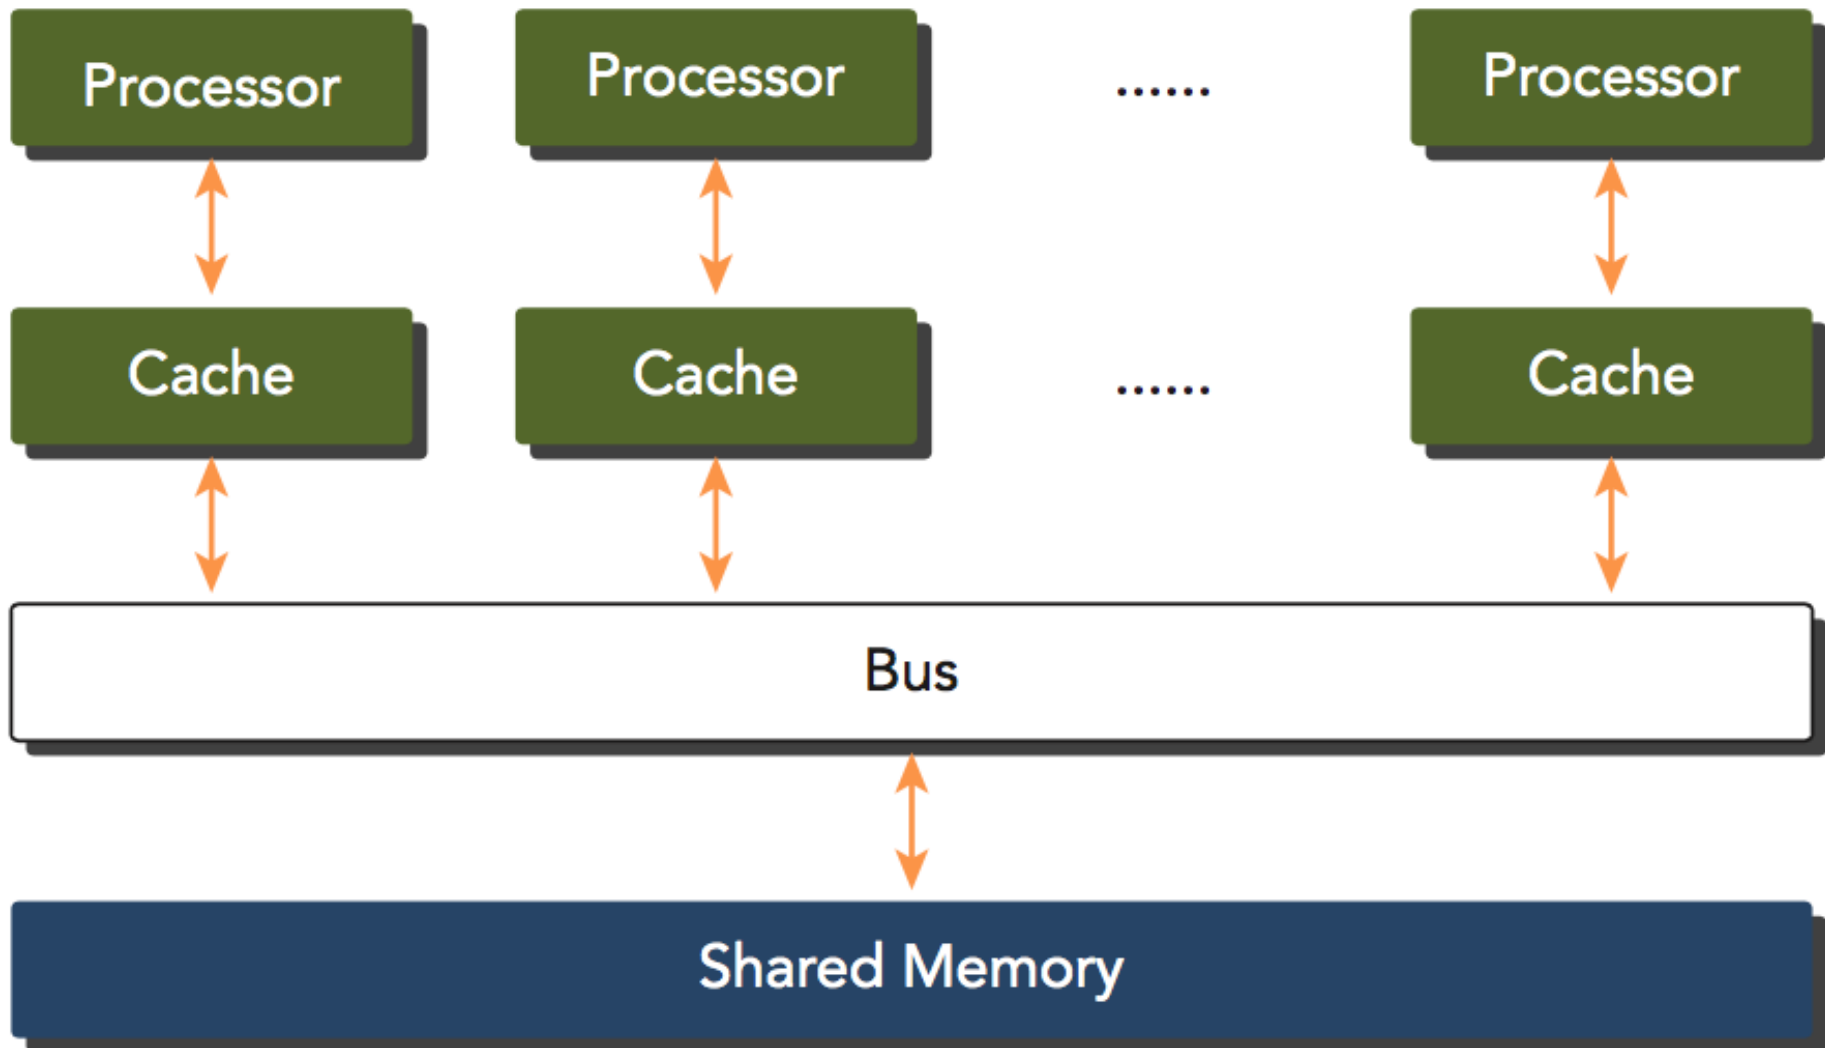
# Architectures parallèles

## Clusters



# Architectures parallèles

Shared memory

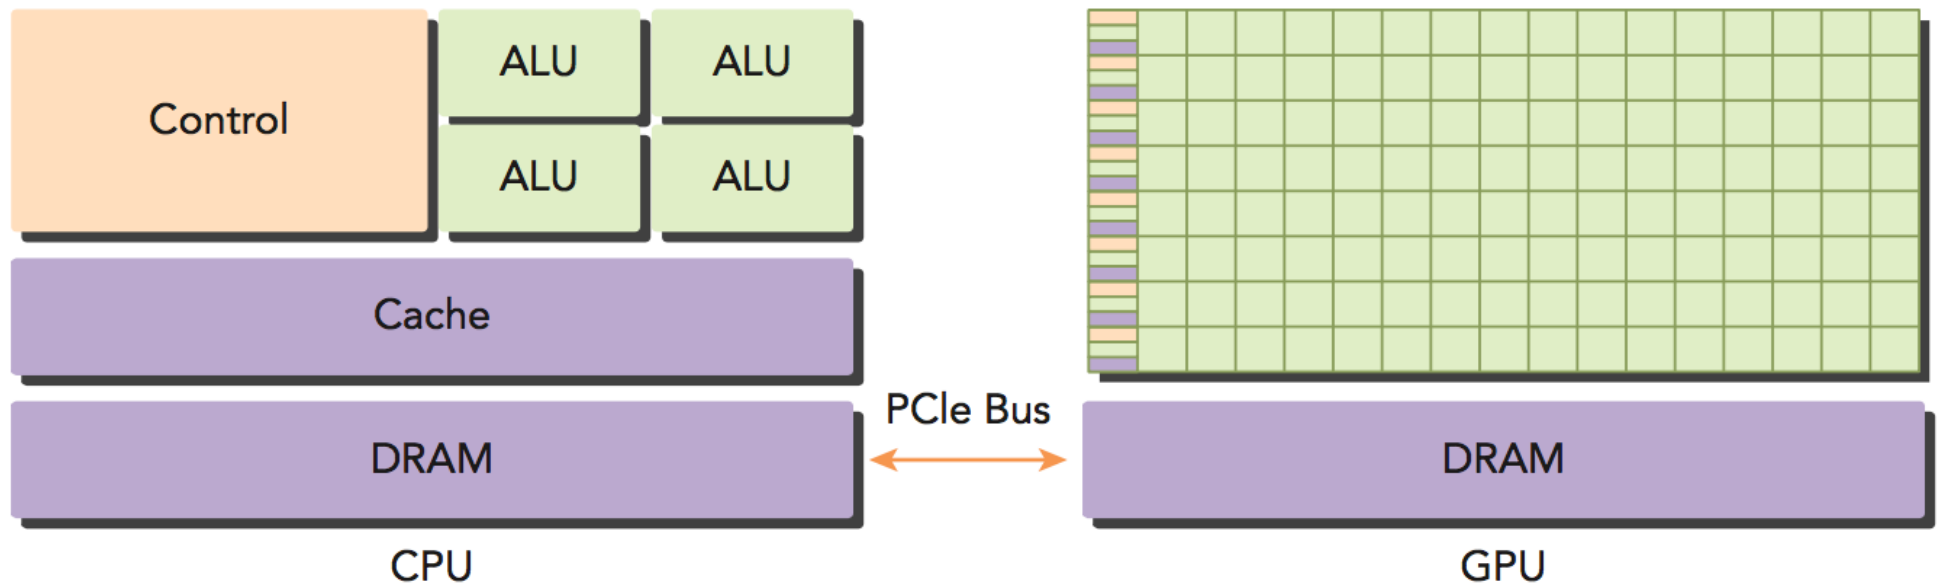


# Architectures parallèles

## Multiprocessors

- Processors can consist of several cores placed on one or several chips
- Many-core : lots of cores (tens or hundreds)
- GPU : many-core. More than SIMD. Called *SIMT*.
  - Threads can have different paths.
  - Threads have their own register state.

# Heterogeneous architecture : master-slave control



# Concepts

- Latency = time to perform an operation (data fetch or computation)
- Throughput = number of operations per sec :  
Gflops, Tflops...
- Bandwidth = data processed per time: Mb/s, Gb/s...

# Parallel computation : software side

- Assign tasks to different threads.
- Two main strategies for handling data : block & cyclic partition.



*Block partition: each thread takes one data block*



*Cyclic partition: each thread takes two data blocks*



# Pros and cons of GPUs

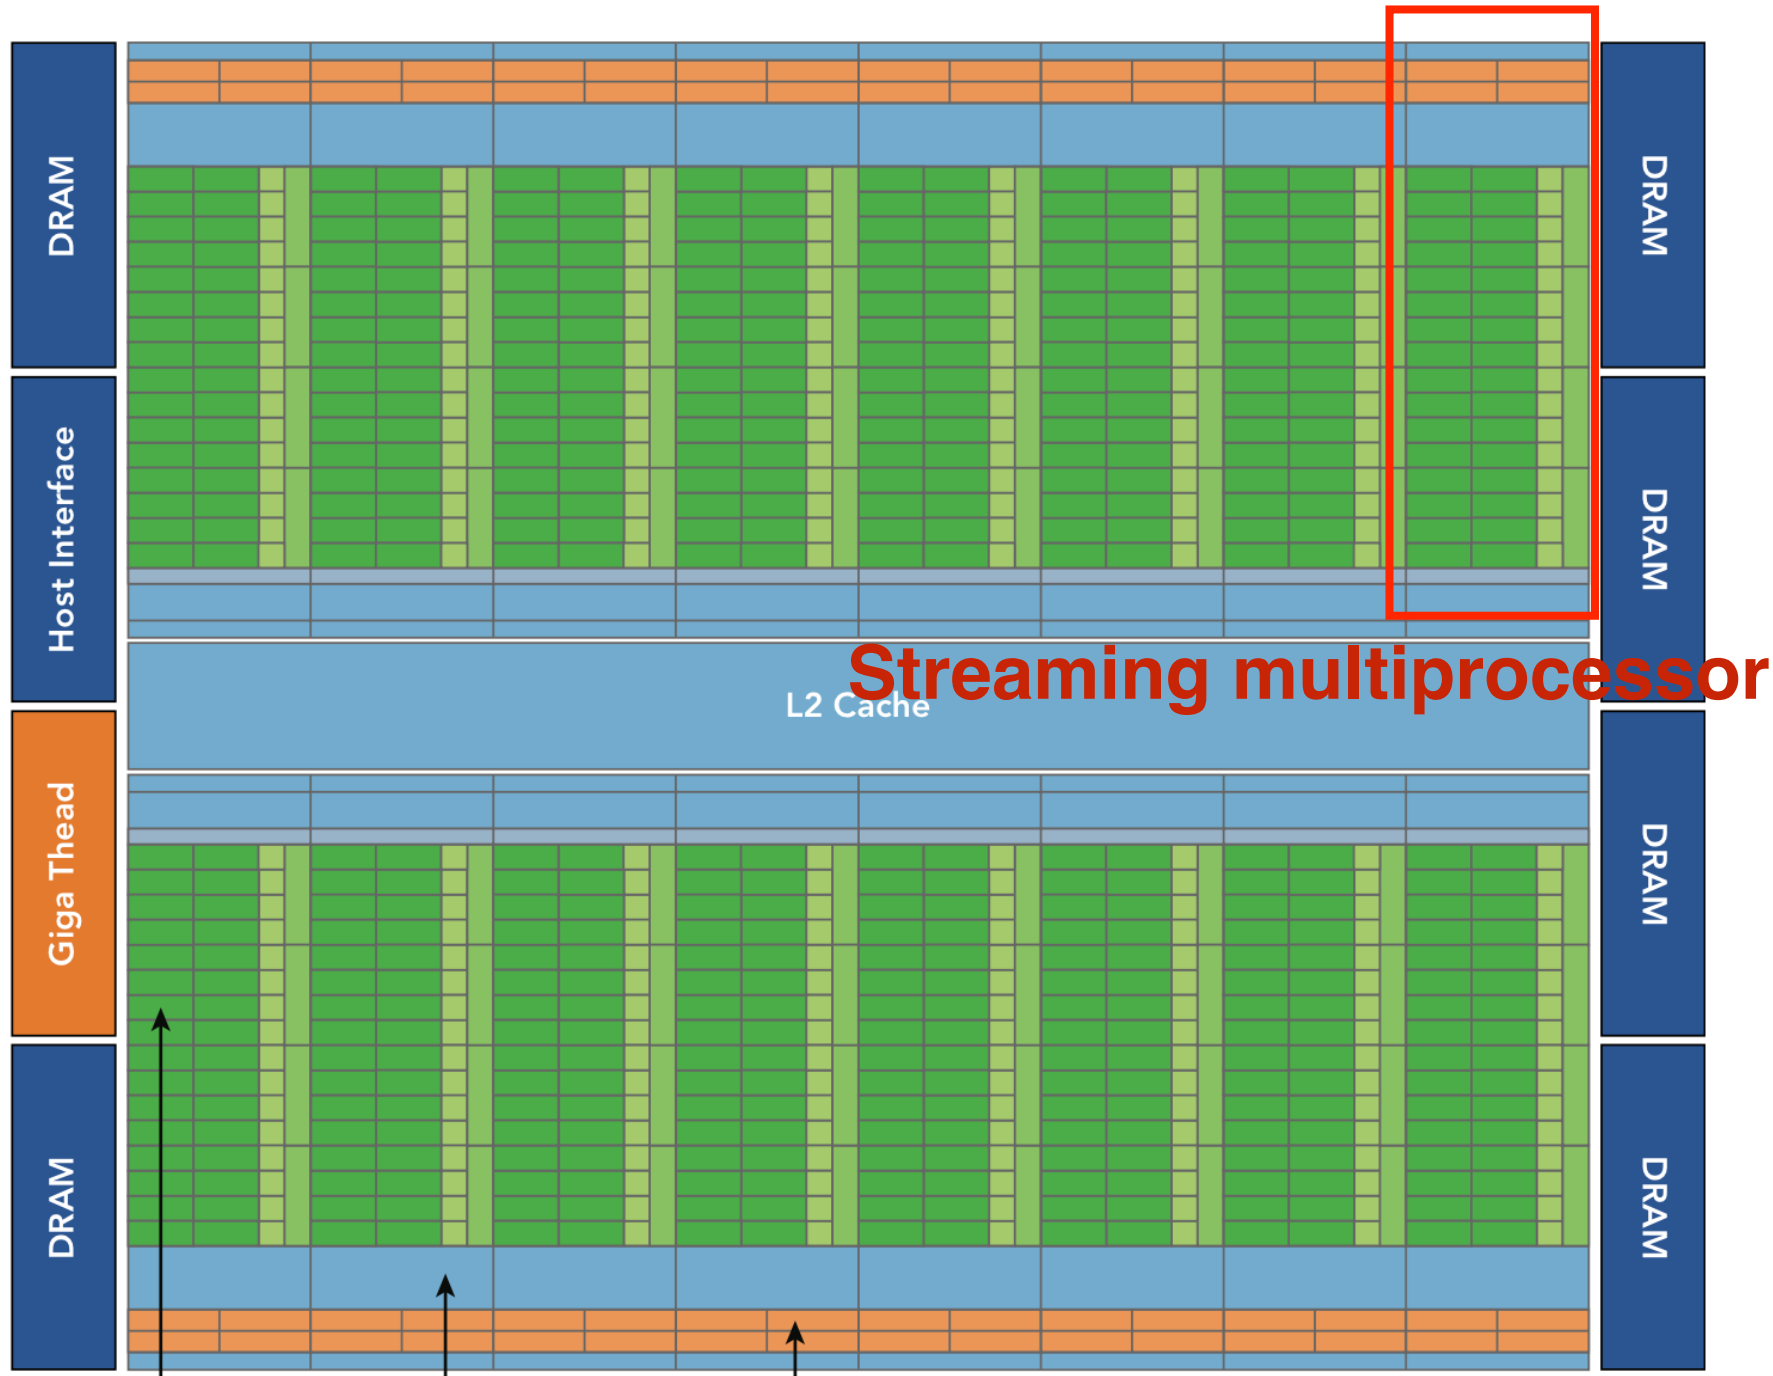
# Pros

- 9 x less space
- 7 x less electrical consumption
- 6 x less expensive

# Cons

- LATENCY !
  - PCI bus slower than frontside bus (RAM)
  - Solutions : Heterogeneous architectures  
(Fusion chez AMD, Sandy-Bridge chez Intel...)

# GPU Architecture





CUDA core

Shared memory,  
register file,  
and L1 cache

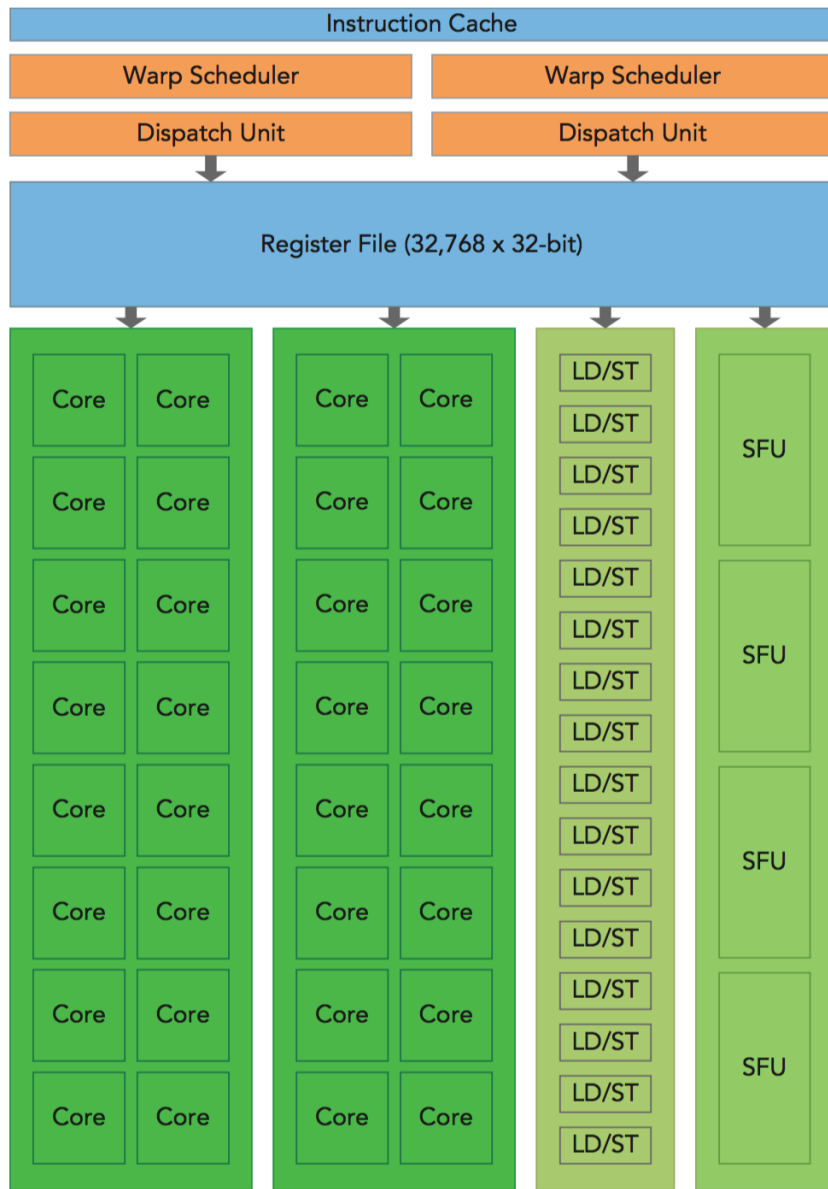
Scheduler and  
dispatch units

Typical GPU

# Streaming multiprocessors



# Streaming multiprocessors

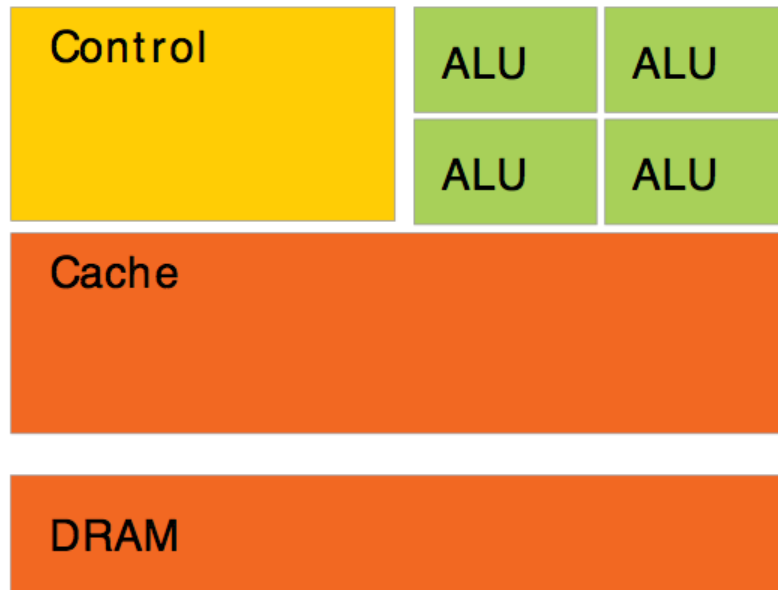


Blocks of threads  
scheduled to one SM


# Streaming multiprocessors

- Minimum « packet » of threads : 32. Called a *wrap*.
- One block threads assigned to one SM -> Memory sharing, threads collaborations...
- You *can* synchronize threads within on block
- You *cannot do it* among blocks





CPU



GPU

More ALUs devoted to computation

# Difference between CPU and GPU threads

- « Heavyweight » vs « lightweight »
- Context switching costless

# Nvidia Families

- Tegra (mobiles and tablets)
- GeForce (multipurpose)
- Quadro (visualisation)
- Tesla (computation)

# Tesla family

- Tesla architecture (attention to confusion !) : 1.0
- Fermi architecture (2.0) : 2010
- Kepler (3.0) : 2012

	FERMI (TESLA C2050)	KEPLER (TESLA K10)
CUDA Cores	448	2 x 1536
Memory	6 GB	8 GB
Peak Performance*	1.03 Tflops	4.58 Tflops
Memory Bandwidth	144 GB/s	320 GB/s

CUDA API

# CUDA API

1.Low-Low level API : driver API (not covered here)

2.Low level : device API

- Mutually exclusive
- No performance difference

# API overview : libraries and packages

- cublas (BLAS) <https://developer.nvidia.com/gpu-accelerated-libraries>
- cublas\_device (BLAS Kernel Interface)
- cuda\_occupancy (Kernel Occupancy Calculation [header file implementation])
- cudadevrt (CUDA Device Runtime)
- cudart (CUDA Runtime)
- cufft (Fast Fourier Transform [FFT])
- cupti (Profiling Tools Interface)
- curand (Random Number Generation)
- cusparse (Sparse Matrix)
- cusolver : combines cuBlas and cuSparse
- npp (NVIDIA Performance Primitives [image and signal processing])
- nvblas ("Drop-in" BLAS)
- nvcuvid (CUDA Video Decoder [Windows, Linux])
- thrust (Parallel Algorithm Library [header file implementation])
- Arrayfire : higher level functions

# Why use libraries ?

- Highly optimized
- Portable and maintained
- Thread safety



# Verification !

Cuda toolkit v 6.0 installers (5.0 is ok also) :  
<https://developer.nvidia.com/cuda-toolkit-60>

- which nvcc (should be /usr/local/cuda/bin on Linux/  
OS X)
- nvcc -V

# MAC

<http://docs.nvidia.com/cuda/cuda-getting-started-guide-for-macos-x/index.html#axzz3YJSArOre>

- Répertoire d'installation :
  - /Developer/NVIDIA/CUDA-xx
  - Symlinks créés dans /usr/local/cuda

# Linux

[http://docs.nvidia.com/cuda/cuda-getting-started-guide-for-linux/  
index.html#axzz3YJSArOre](http://docs.nvidia.com/cuda/cuda-getting-started-guide-for-linux/index.html#axzz3YJSArOre)

- Répertoire d'installation :
  - /usr/local/cuda

# Examples coming with cuda toolkit

- Go to `/usr/local/cuda/samples`
- Make for all samples (can take some time)
- Or make inside a specific application

Hardware exploration

# deviceQuery

- Go in  
/usr/local/cuda/samples/1\_Uutilities/deviceQuery
- Run make and execute
- Sample outputs next slides

Device 0: "GeForce GT 650M"

CUDA Driver Version / Runtime Version	5.5 / 5.5
CUDA Capability Major/Minor version number:	3.0
Total amount of global memory:	1024 MBytes (1073414144 bytes)
( 2) Multiprocessors, (192) CUDA Cores/MP:	384 CUDA Cores
GPU Clock rate:	900 MHz (0.90 GHz)
Memory Clock rate:	2508 Mhz
Memory Bus Width:	128-bit
L2 Cache Size:	262144 bytes
Maximum Texture Dimension Size (x,y,z)	1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
Maximum Layered 1D Texture Size, (num) layers	1D=(16384), 2048 layers
Maximum Layered 2D Texture Size, (num) layers	2D=(16384, 16384), 2048 layers
Total amount of constant memory:	65536 bytes
Total amount of shared memory per block:	49152 bytes
Total number of registers available per block:	65536
Warp size:	32
Maximum number of threads per multiprocessor:	2048
Maximum number of threads per block:	1024
Max dimension size of a thread block (x,y,z):	(1024, 1024, 64)
Max dimension size of a grid size (x,y,z):	(2147483647, 65535, 65535)
Maximum memory pitch:	2147483647 bytes
Texture alignment:	512 bytes
Concurrent copy and kernel execution:	Yes with 1 copy engine(s)
Run time limit on kernels:	Yes
Integrated GPU sharing Host Memory:	No
Support host page-locked memory mapping:	Yes
Alignment requirement for Surfaces:	Yes
Device has ECC support:	Disabled
Device supports Unified Addressing (UVA):	Yes
Device PCI Bus ID / PCI location ID:	1 / 0
Compute Mode:	

< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 5.5, CUDA Runtime Version = 5.5, NumDevs = 1, Device0 = GeForce GT 650M

Result = PASS



Device 1: "Tesla M2050"

CUDA Driver Version / Runtime Version 5.0 / 5.0  
CUDA Capability Major/Minor version number: 2.0  
Total amount of global memory: 3072 MBytes (3220897792 bytes)  
(14) Multiprocessors, ( 32) CUDA Cores/MP: 448 CUDA Cores  
GPU Clock rate: 1147 MHz (1.15 GHz)  
Memory Clock rate: 1546 Mhz  
Memory Bus Width: 384-bit  
L2 Cache Size: 786432 bytes  
Maximum Texture Dimension Size (x,y,z) 1D=(65536), 2D=(65536, 65535), 3D=(2048, 2048, 2048)  
Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers  
Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers  
Total amount of constant memory: 65536 bytes  
Total amount of shared memory per block: 49152 bytes  
Total number of registers available per block: 32768  
Warp size: 32  
Maximum number of threads per multiprocessor: 1536  
Maximum number of threads per block: 1024  
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)  
Max dimension size of a grid size (x,y,z): (65535, 65535, 65535)  
Maximum memory pitch: 2147483647 bytes  
Texture alignment: 512 bytes  
Concurrent copy and kernel execution: Yes with 2 copy engine(s)  
Run time limit on kernels: No  
Integrated GPU sharing Host Memory: No  
Support host page-locked memory mapping: Yes  
Alignment requirement for Surfaces: Yes  
Device has ECC support: Disabled  
Device supports Unified Addressing (UVA): Yes  
Device PCI Bus ID / PCI location ID: 131 / 0  
Compute Mode:

< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

> Peer access from Tesla M2050 (GPU0) -> Tesla M2050 (GPU1) : No

> Peer access from Tesla M2050 (GPU1) -> Tesla M2050 (GPU0) : No

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 5.0, CUDA Runtime Version = 5.0, NumDevs = 2, Device0 = Tesla M2050, Device1 = Tesla M2050

# Bandwidth test

- Run the example in  
`/usr/local/cuda/samples/1_Uutilities/bandwidthTest`

# NVIDIA Management Library (NVML)

<https://developer.nvidia.com/nvidia-management-library-nvml>

[http://developer.download.nvidia.com/compute/cuda/6\\_0/rel/gdk/nvidia-smi.331.38.pdf](http://developer.download.nvidia.com/compute/cuda/6_0/rel/gdk/nvidia-smi.331.38.pdf)

[https://www.microway.com/hpc-tech-tips/nvidia-smi\\_control-your-gpus/](https://www.microway.com/hpc-tech-tips/nvidia-smi_control-your-gpus/)

- C-based API useful for :
- Monitoring GPU use
- Modifying some settings
- Available on Linux and Windows

- Command line tool : nvidia-smi
- In c code : `#include <nvml.h>`