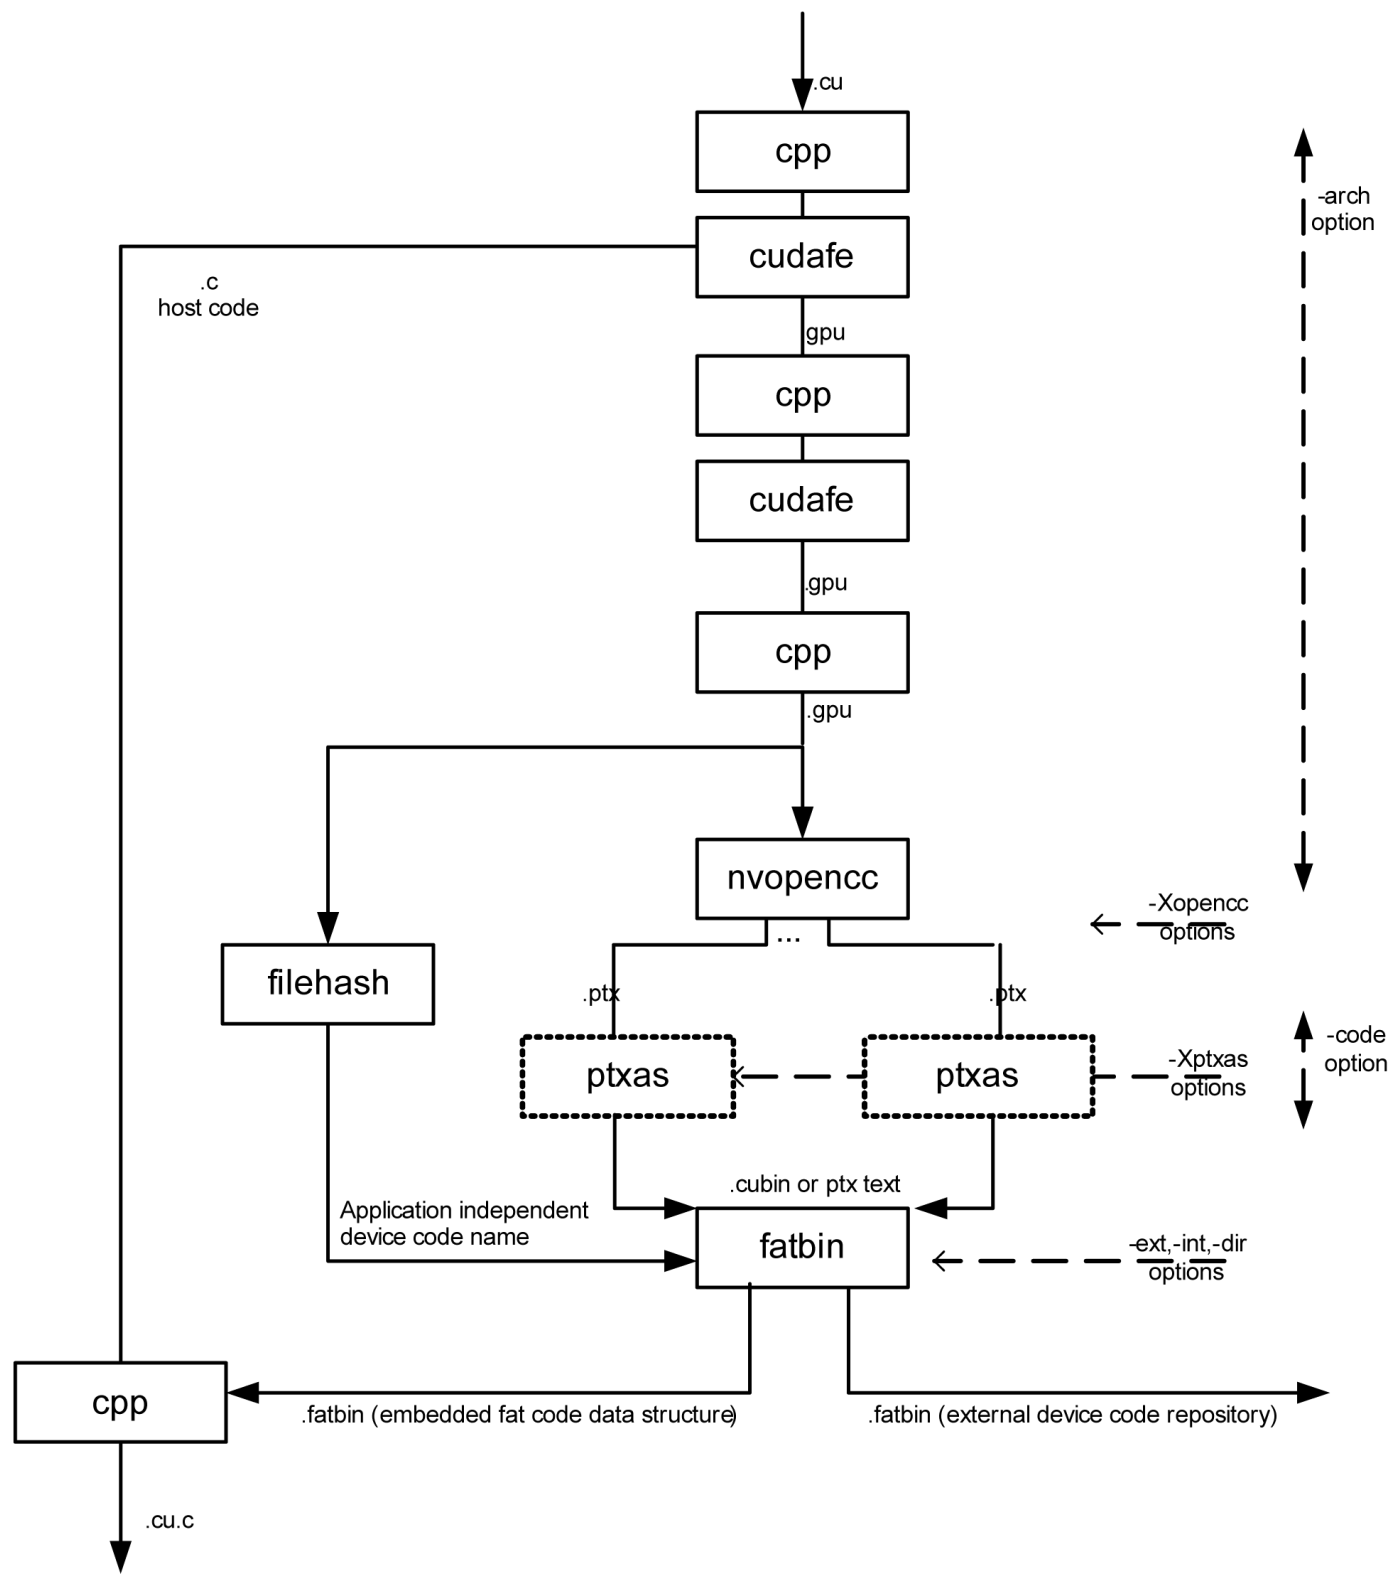


Compiler et pipeline

- Compilateur nvidia : nvcc
 - Calls host compiler c/c++ for c/c++ code et ptx for cuda code

Compilation phases overview (*trajectory*)



- cudafe (front end) : splits the code in .gpu (device code) and c code
- After : .gpu -> .cubin and / or .ptxas
- Flows through a *descriptor*
- Finally, descriptor included into *host code*
- When program is running => descriptor inspected by CUDA runtime.

With what to feed it ?

- Source files
 - .cu : CUDA source
 - .cpp / .cc / .cxx
- Des objets :
 - .o / .obj / .so
- Des librairies :
 - .lib / .a

But also...

- Intermediate assembly : .ptx
- Ressource : .res
- Gpu intermediate file : .gpu
- Pre-processed : .cup

Note : only compile .c files
with nvcc if only C API is
present ! (no kernels)

Step 0 : warming up

- Write a dummy C main function in a file with .c extension (eg : hello world)

Case 1 : compile a c file with only API functions

- Write any of the API functions. Example :
float *ptr;
cudaMalloc(&ptr, 1 * sizeof(float));
- Compile with nvcc. Works ?

Case 2 : compile a C file with a kernel inside

- Now place this block visible by the main (just at the top of main function or in a header file)
`__global__ void foo();`
- Compile. Works?

Case 2 : compile a C file with a kernel inside

- Try again by changing the extension from .c to .cu

Case 3 : you want to be able compiling with gcc (no gpu)

- Remove the API function form inside the main (i.e. you can leave an empty main) but let the `__global__` function.
- Compile with gcc (or any c compiler).

Case 3 : you want to be able compiling with gcc (no gpu)

- You want to keep your source code, that can run on non gpu machines
- Try wrapping the `__global__` function by checking the definition of macro `__NVCC__` :

```
#ifdef __NVCC__  
__global__ void foo();  
#endif
```

Case 4 : you want to have your GPU functions in a separate file from the c file but keep the .c extension for the main

- You CANNOT have a GPU function (`__global__`, `__device__`...) inside a c / cpp file !
- Solution : compile files separately and link with a host c compiler. Don't forget to link with cuda runtime library (`-L/usr/local/cuda/lib -lcudart`)
- You have to use a C function (wrapper) to call your global function with the command
`foo<<<1,1>>>(void);`
- Solution at next slide

Case 4 : you want to have your GPU functions in a separate file from the c file but keep the .c extension for the main

- Consider the files main.c and side.cu in COMPIRATION folder.

1. gcc -c main.c # generates a object file

2. nvcc -c side.cu # also

3. gcc main.o side.o -L/usr/local/cuda/lib -lcudart

Case 5 : you want to dispatch your device functions into multiple files

- Function executing from the *device*, i.e. from a `__global__` function, example :
`__device__ int dev();`
- Place this declaration in `side.cu` and call it (like you would with any function) from the `__global__` function in `side.cu`
- Compile it. What happens?

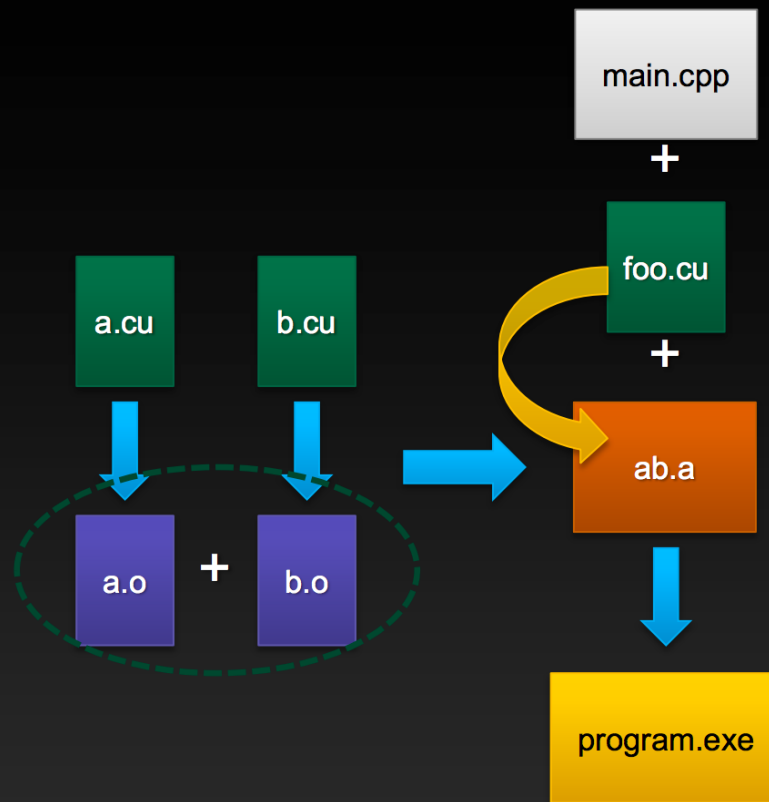
« Separate compilation »

aka « relocatable device code » (RDC)

Starting from 5.0 !

Separate compilation

CUDA 5: Library Support

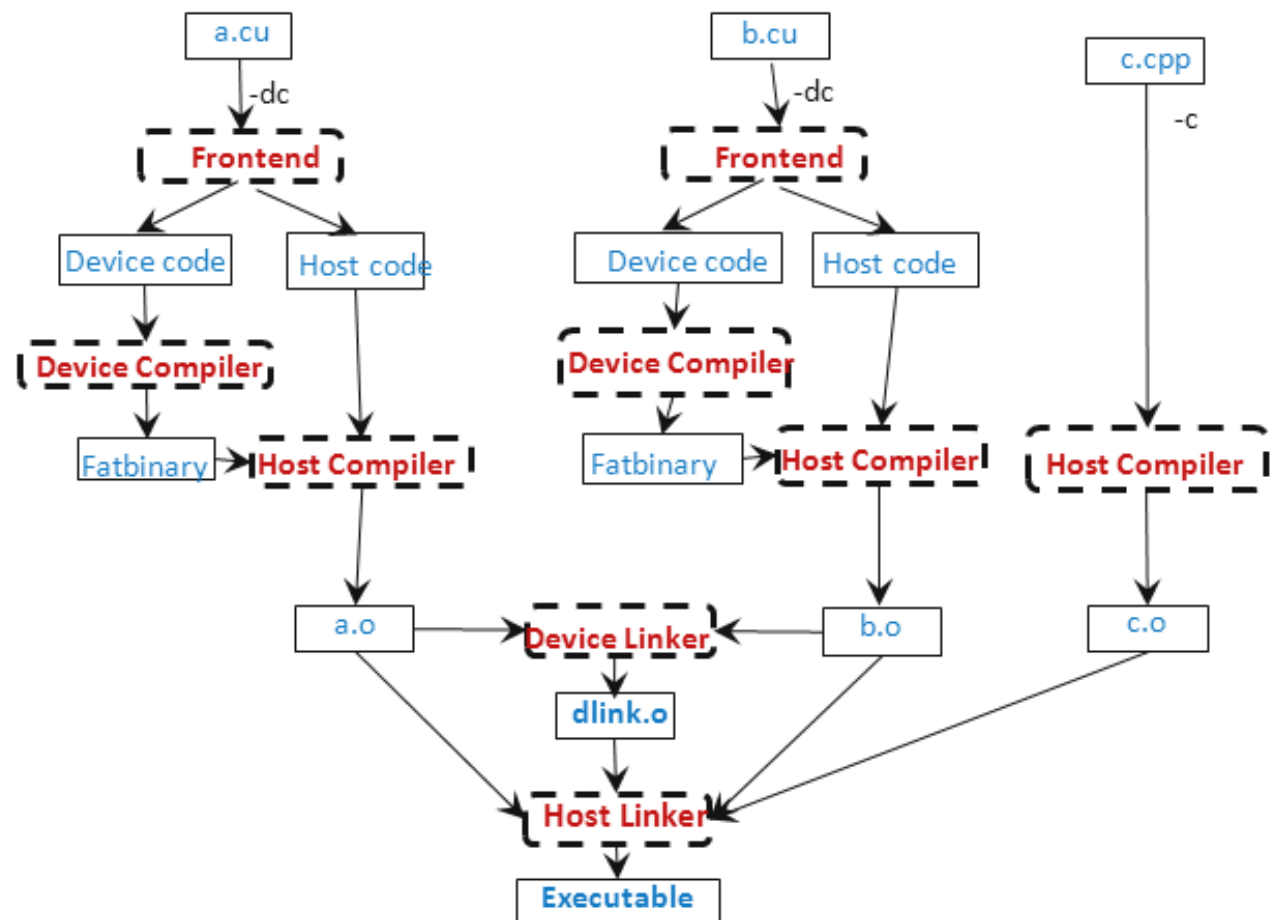


Can combine object files into static libraries

Link and externally call *device* code

- Avant : tout le *device code* dans un seul fichier !
- -dc : embed du *relocatable device code* dans du *host code* qui sera lié par *nvlink* (device linker) lors du linkage.

- nvlink peut être appelé explicitement par -dlink
- utilisation possible de *extern*



Case 6 : example of architecture code selection

- Write a .cu file with a `__global__` function doing a `printf(« hello world! \n»);`
- Call it from the main with `foo<<<1, 1>>>();`
- Can you compile ?

Case 6 : example of architecture code selection

- Solution in `hello_world.cu` in the `COMPILATION` folder.
- `printf` is only available for architecture `>= 2.0`.
- Wrap the function with a `#if (__CUDA_ARCH__ >= 200)`

Compiler flags /
options

Les plus communs

- `-o file` : le nom du output file
- `-c file` : compile le fichier en fichier objet
- `-I path` : inclut un chemin de recherche pour les headers
- `-L path` : inclut un chemin de recherche pour les librairies (phase de linkage)
- `-O level` : niveau d'optimisation (1, 2, 3...)
- `-clean` : supprime les fichiers temporaires qui auraient du être créés

Suite

- -m *arch* : machine (32 ou 64)
- -g : génère du code débuggable
- -G : génère du code device débuggable
- -shared : génère une librairie partagée
- -M *file* : générer les dépendances pour utiliser en Makefile
- -D *macro* : définit une macro
- -U *macro* : dédéfinit une macro (undefine)

Parlons architecture

Architecture hardware / real (ex : sm_20) + architecture virtuelle / compute (ex : compute_20)

Architectures GPUs

- Tesla
 - sm_11
 - sm_12
 - sm_13
- Fermi
 - sm_20
 - sm_21
- Kepler
 - sm_30
 - sm_32
 - sm_35
- Maxwell
 - sm_50
 - sm_52

Rien d'assuré pour la compatibilité des fichiers binaires entre versions!

Good to know : compilation macros

- 3 preprocessor macros are defined by nvcc :
 - `__NVCC__` : When nvcc is used to compile the file
 - `__CUDACC__` : When you compile cuda code (but not c / cpp)
 - `__CUDACC_RDC__` : When you compile cuda code with « relocatable device code » (RDC)

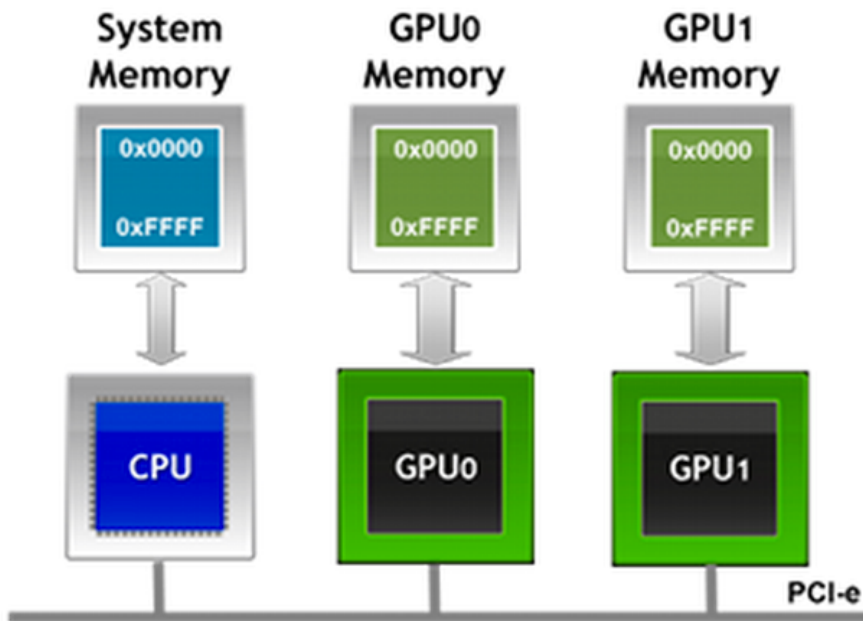
Versions : What they brought to us

Release notes

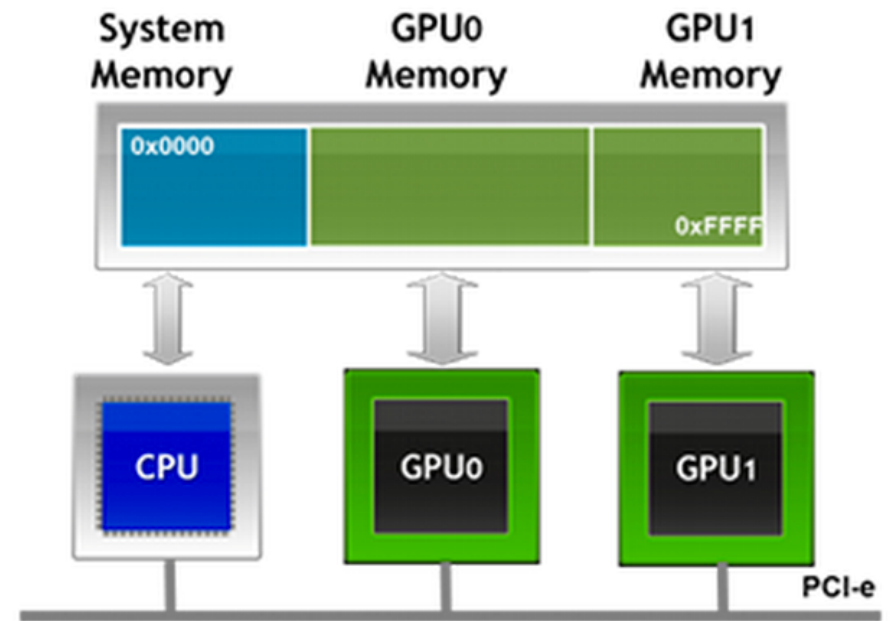
- 4.0 (may 2011) unified virtual memory addressing
- 5.0 (oct. 2012)
 - dynamic parallelism (with Capability 3.5)
 - separate compilation (<http://on-demand.gputechconf.com/cexpress/2012/presentations/gpu-object-linking.pdf>)
- 6.0 (april 2014)
 - transparent memory copies handling (utilisation d'un pointe unique)
 - XT libraries (cublas & cufft) for automatic scaling to multiple GPUs (!)
 - Support for Maxwell architecture (sm_5.0)
- 7.0 (march 2015)
 - C++11 support !
 - cuSOLVER library (linear operations for eigen problems, de and sparse solvers)

Unified virtual memory addressing (> 4.0)

No UVA: Multiple Memory Spaces



UVA: Single Address Space



- +
- added routines to CUBLAS & CUSPARSE
- added math functions (sincos, sincospi, cospi, sinpi...) and performance improvements
- added distributions and generators to CURAND

Gcc/G++ compatible versions

5.0 : up to gcc 4.7

6.0 : up to gcc 4.8

7.0 : up to gcc 4.9

Note : C++ 2011 only supported starting from 7.0 !