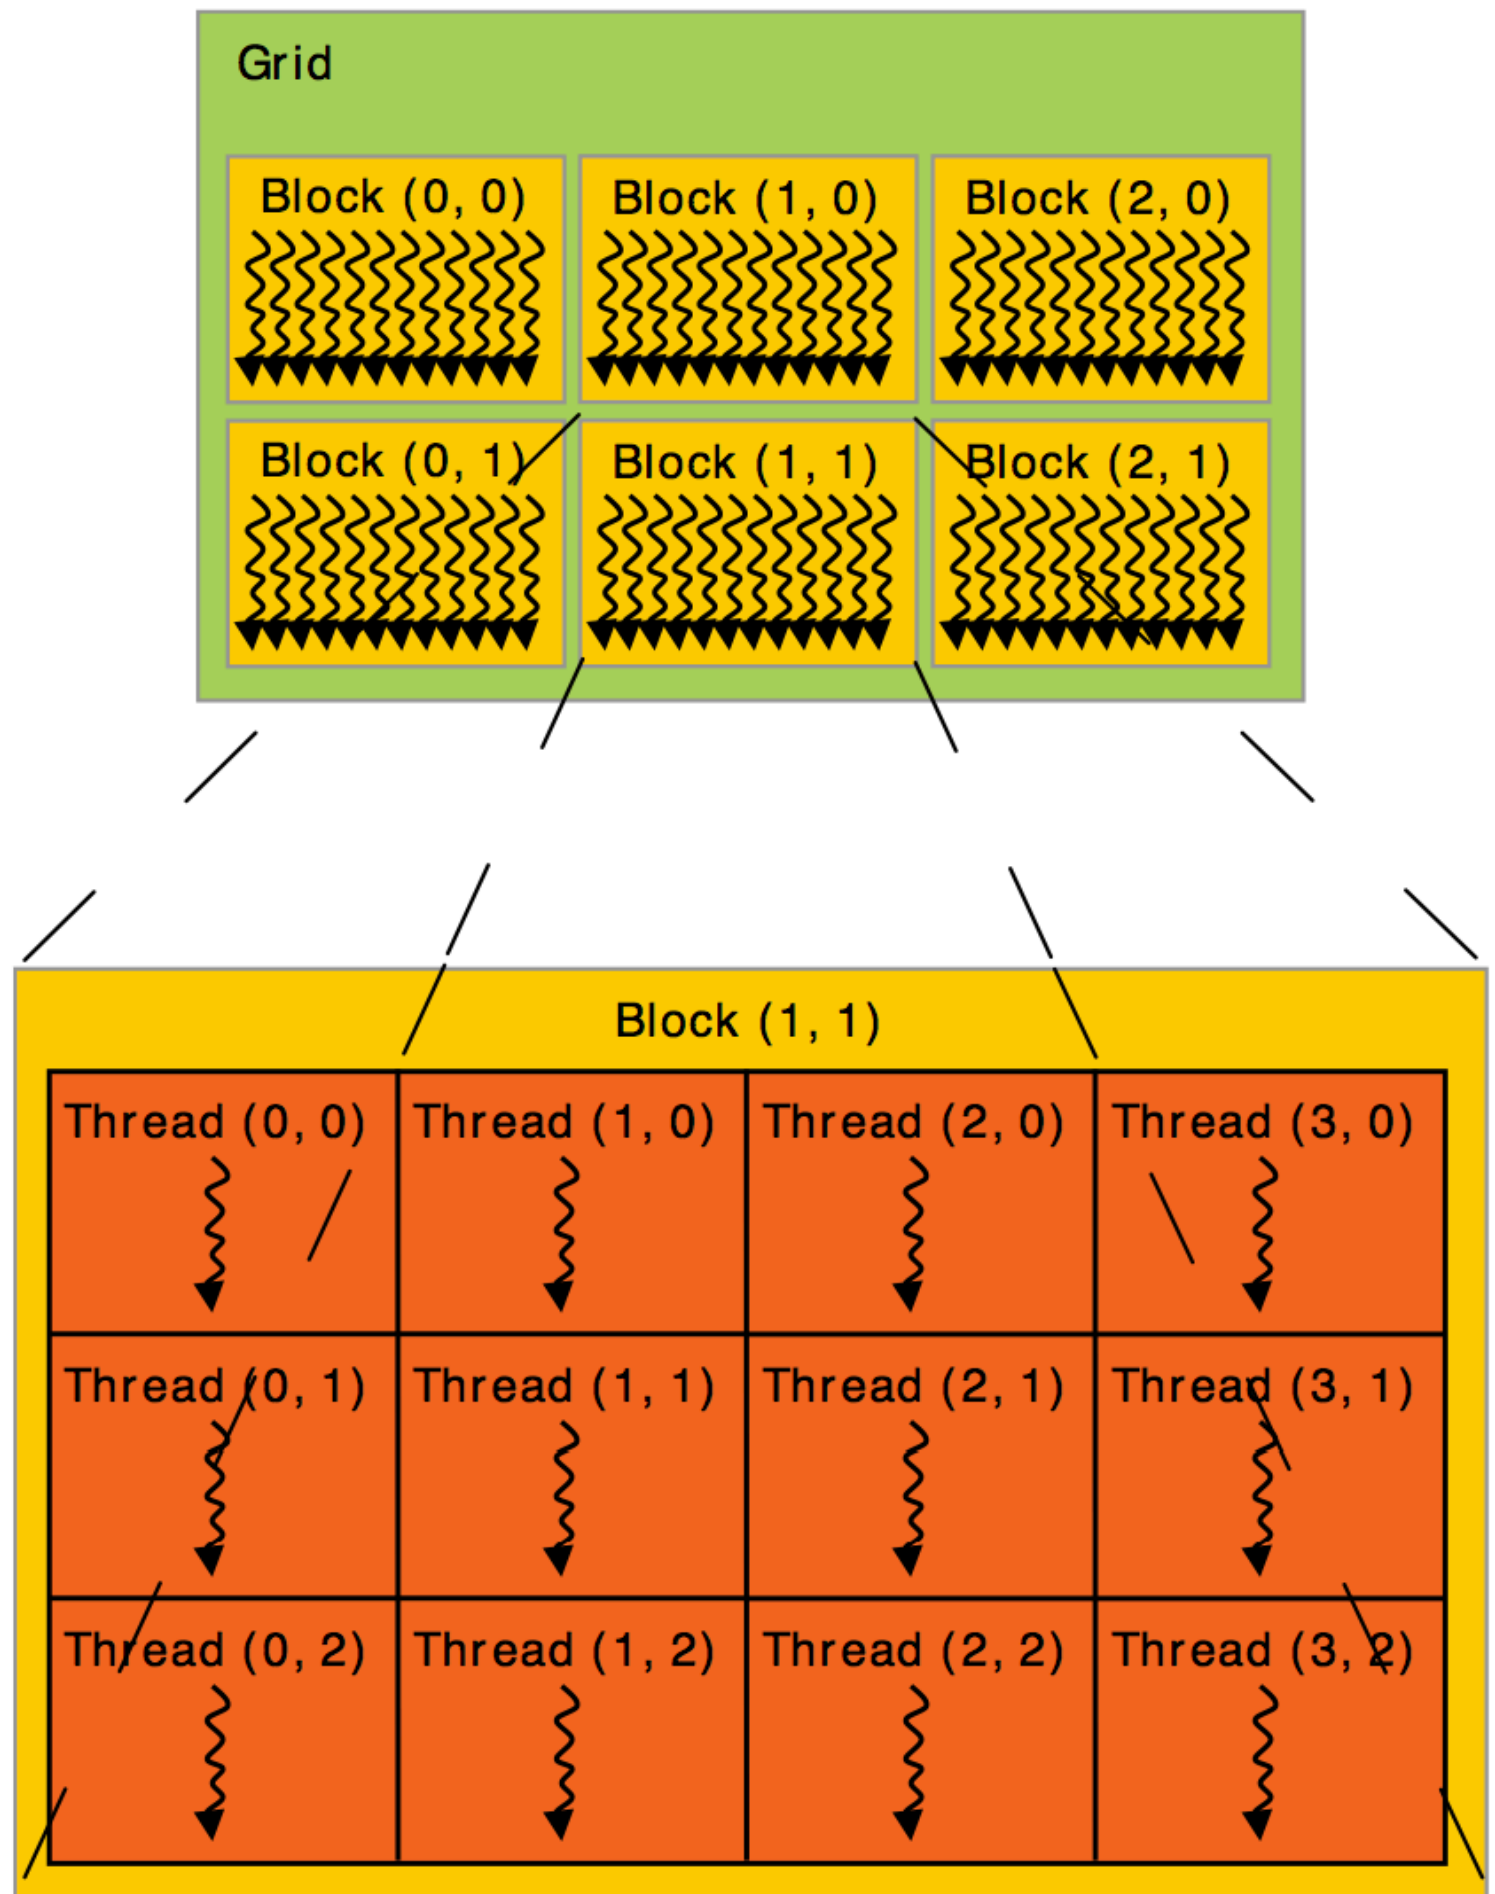


Writing your own
device functions

- Kinds of functions :
 - * The one that you call from host (CPU) are called *kernels* and have prefix `__global__`
 - ❖ Always return void
 - * The ones that you call from kernels : `__device__`
 - * The general ones that you could call both from device and host : `__host__ __device__`
 - * Except for dynamic parallelism, kernel cannot launch a kernel

- Example of a kernel
 - * `__global__ void myKernel (int arg1, float *arg2) (no ellipsis)`
- Launch with
 - * `myKernel<<<nBlocks, nThreadsPerBlock>>>(arg1, arg2, ...)`
 - * Use `dim3` type for grid dimension and block dimension for 2D or 3D blocks
- Know the index of the working thread!
 - * Use `threadIdx.x`, `blockIdx.x`, `blockDim.x`, `gridDim.x` (also `y` and `z`)

Example : two dimensional grid of two dimensional blocks



Hello World

- KERNELS / [hello_world.cu](https://github.com/kerneLS/hello_world.cu)

Computes operations on an array

- Allocate memory on device if not already done
 - * use `cudaMalloc(&pointer, memSize)` on a *pointer to pointer* to alloc heap memory on device
- `cudaMemSet(pointer, value, memSize)`

- Array source :
 - * CPU vector copied to device
 - ❖ with cudaMemcpy(destination_pointer, source_pointer, memSize, direction_of_the_copy)
 - ❖ Utilisation « easier » since 6.0
 - * thrust raw_pointer

Exercice 1

- Create two sequences on the CPU, S1 ranging from 1 to 100 with step 1, S2 ranging from 1 to 100 with step 2
- Transfer the arrays to device
- Compute the sum of both arrays in a 3rd vector, using a blockSize of 32
- Don't forget CUDA_CHECK and CUDA_CHECK_ERROR !