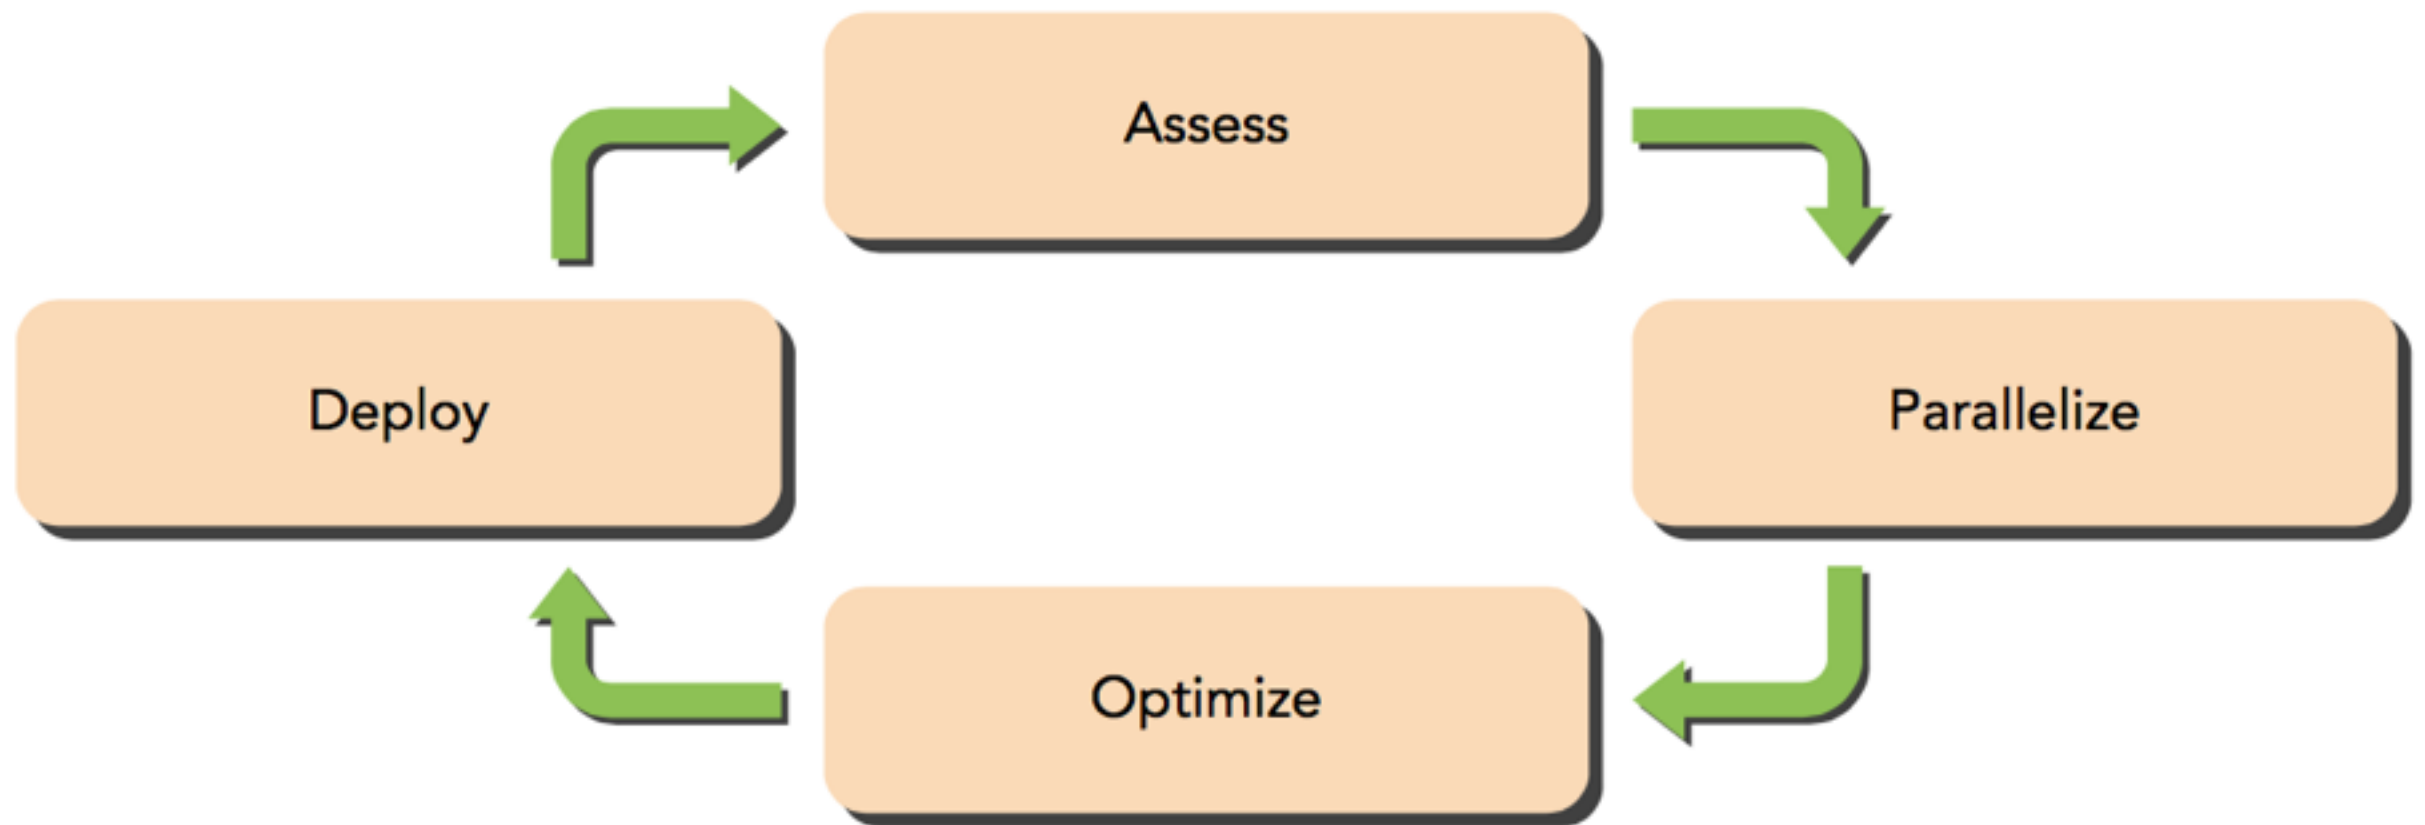


OPTIMIZATIONS

Memory components, loop unrolling, thread shuffle, ninja functions, memory access...

General *strategy*



APOD

1. Assessment

- Identify bottlenecks
- High priority : parallelizable loop structures with lots of computation. Even in already host-parallelized sections.

2. Parallelize

- Cuda libraries
- Kernels
- OpenACC

3. Optimization

3.1 Grid Level

- Concurrency, multiple kernels at the same time
- Overlapping computation and data transfers...

3. Optimization

3.2 Kernel Level

- Memory coalescence & efficient access patterns
- Reduce memory access latency
- Reduce wrap divergence

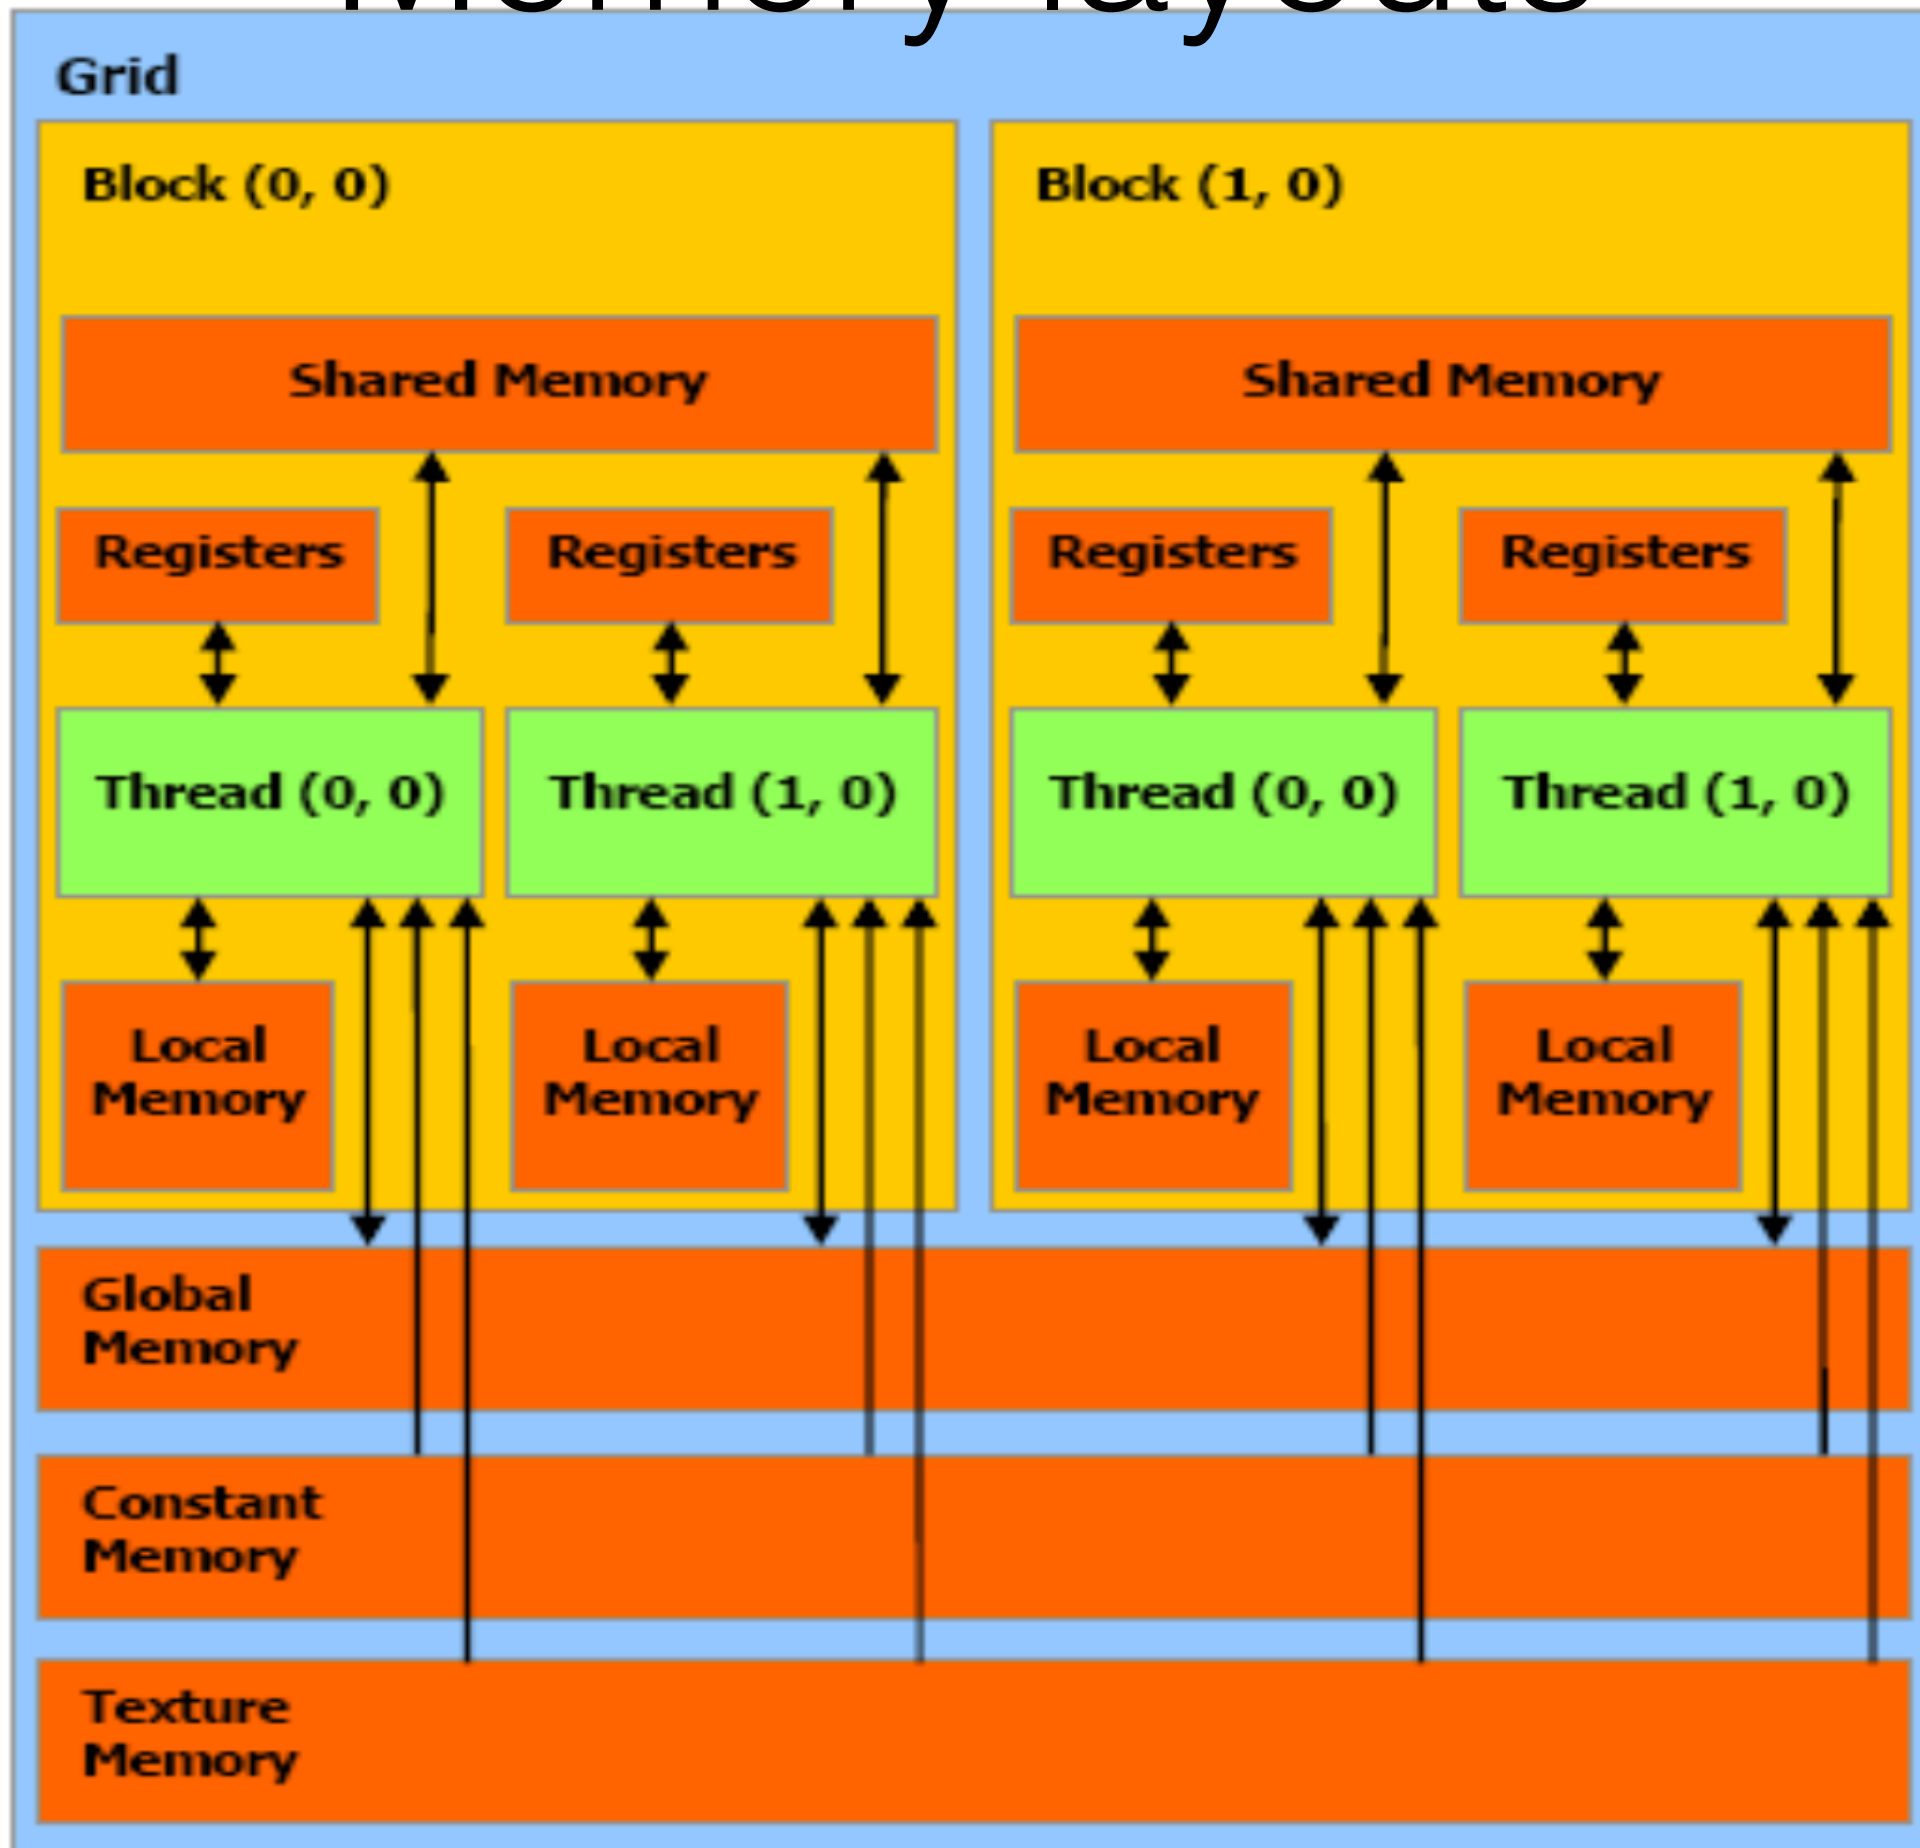
4. Deployment

- Ensure it can run even on non CUDA-capable GPU environnement
- Iterate once again

3.2 Kernel optimization

- Two further optimisations directions : is you kernel :
 - Memory bound ? (most of the time)
 - Computation bound ?

Memory layouts



- Preliminary notes :
- Test in folder OPTIMIZATIONS
- All tests, profiling... should be done in calculation platform
- Using nvprof on cards use for desktop is not accurate

3.2.1 Memory bound kernels

- Goals
- Improve memory access (bytes / sec on the bus)
- Improve concurrent memory access
- Hide latency
- NOTE : memory requests are issued per wrap !

Types of readings

- Cached
- Uncached
- Read-only
- Examples in OPTIMIZATION folder


3.2.1 Memory bound kernels

Load granularity

- Cached : 128 bytes cache-line
- Uncached & read-only : 32 bytes
- Can be chosen at compile time :
 - -Xptxas -dlcm=cg : disabled
 - -Xptxas -dlcm=ca : enabled

Load granularity : L1-cached

- Best memory transactions : *aligned & coalesced*. Better with 32 times (a wrap) 4 bytes (= one 128bytes cache-line)




Aligned & coalesced reading.

If not, several 128-bytes cache lines must be read : *wasted* bandwidth.


Load granularity : L1-cached

- Still ok if order is shuffled : line is cached, only one transaction needed, bus use = 100%




Load granularity : L1- cached

- Misaligned reading : two lines read. bus use = 50%




Load granularity : L1-cached

- Only one element read : two lines read. bus use = $1/32 \sim 3.1\%$




Load granularity : L1-cached

- Very bad access : up to 32 lines read



Load granularity : L1-cached

- Very bad access : up to 32 x 128bytes-lines read




Load granularity : L1- cached

- Note : difference between CPU and GPU L1 cache : no temporal locality (frequent access does not imply data staying in cache)


Load granularity : uncached

- Coalesced & aligned access for 32 bytes lines




Load granularity : uncached

- Shuffled acces : ok




Load granularity : uncached

- Mis-aligned : bus use at least 80 %. Better than with L1-cache




Load granularity : uncached

- Only one element read : bus use = 12.5 %. 4 x better than with L1 cache



Load granularity : uncached

- Very bad access : up to 32 x 32 bytes read, still better



Read-only cache : 32 bytes

- After CC 3.5 : can choose to use read-only with :
 - `function _ldg`
 - `const __restrict__` on pointers

readSegment.cu

- usage : readSegment <offset>
- Try different offsets. Do you see a difference ?
- Use nvprof —metrics gld_efficiency readSegment <offset>

- Try the same thing by disabling the L1 cache
- Compile with the flag `-Xptxas -dlcm=cg`

3.2.1 Memory bound kernels

writeSegment.cu


- Stores are performed at a 32-byte segment granularity

- `nvprof —metrics gld_efficiency —metrics
gst_efficiency ./writeSegment`

3.2.3 AoS vs SoA

`./simpleMathAoS` and `simpleMathSoA` examples

Array of structures (AoS) vs Structure of arrays (SoA)



./simpleMathAoS and simpleMathSoA examples

- Compare the load efficiencies between the two

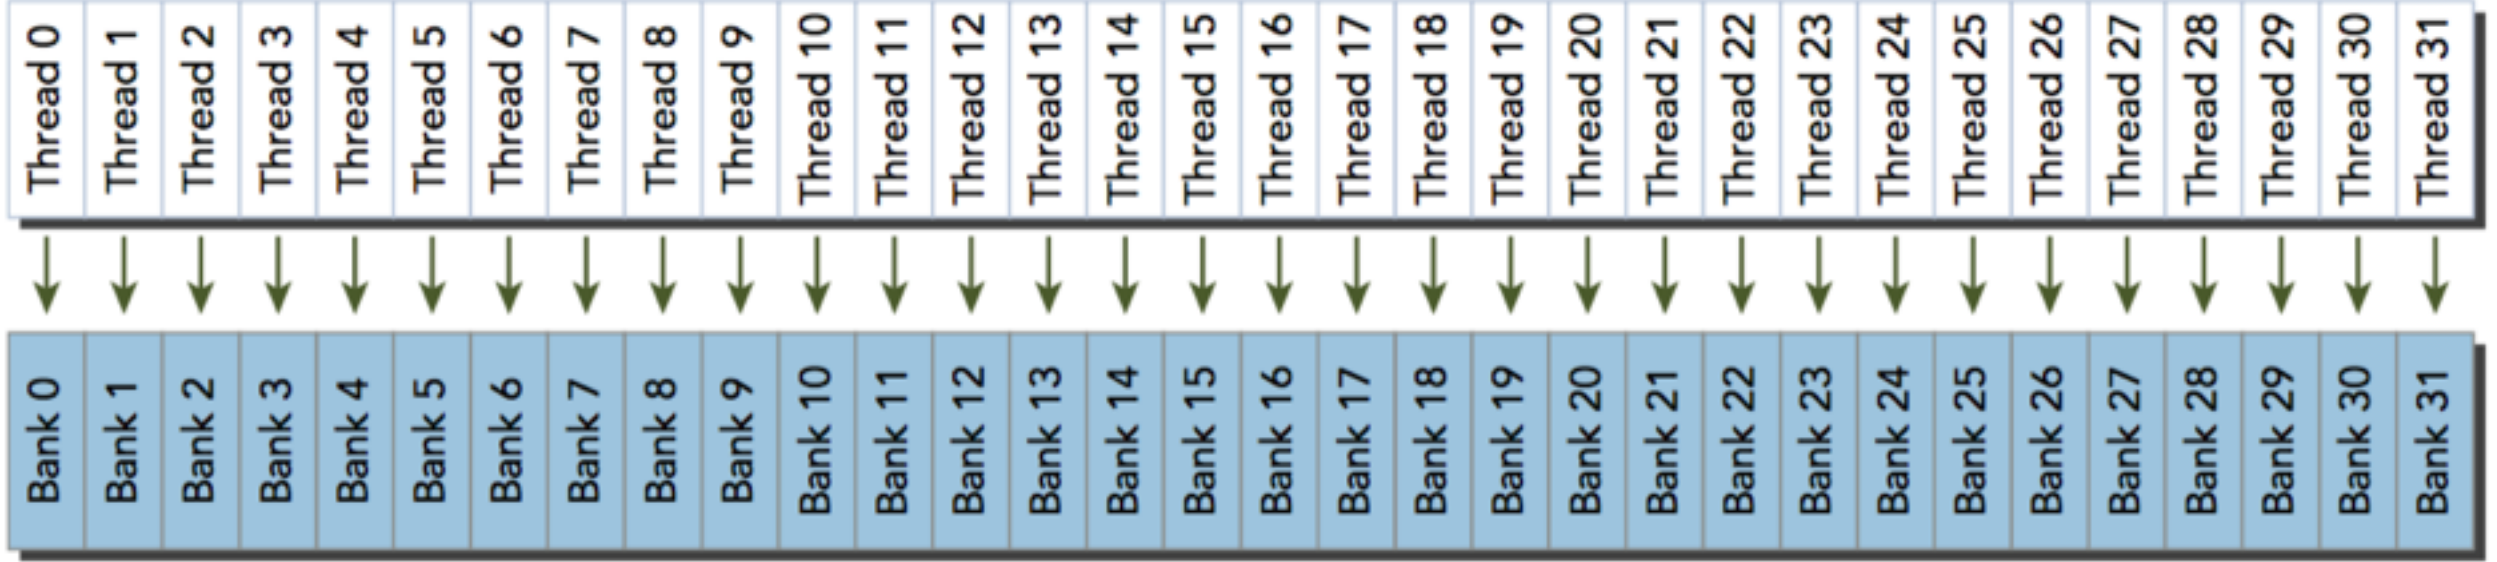
3.2.4 Shared memory

- On-chip, superfast : good for reuse. Shared *per-block*.
- Organised in *banks* of either 4 or 8 bytes (you can choose, will affect performance).
- Access of different words from same bank : *bank conflict*.
- Solution : spread memory, pad arrays.


Types of access

- Parallel access: multiple addresses accessed across multiple banks
- Serial access: multiple addresses accessed within the same bank
- Broadcast access: a single address read in a single bank

Parallel access



Bank conflict : multiplies number of serial transactions
No conflict for threads in the same warp



Allocations

- Inside a kernel :
 - `__shared__ float array[N];`

Configuration

- Choose for the whole device
- Choose also per kernel
- Depends on your kernel : use more share memory or more registers ?

Configuration

- For the whole application :
`cudaError_t cudaDeviceSetCacheConfig(cudaFuncCache cacheConfig);`
 - `cudaFuncCachePreferNone`: no preference(default)
 - `cudaFuncCachePreferShared`: prefer 48KB shared memory and 16 KB L1 cache
 - `cudaFuncCachePreferL1`: prefer 48KB L1 cache and 16 KB shared memory
 - `cudaFuncCachePreferEqual`: prefer 32KB L1 cache and 32 KB shared memory

Configuration

- For a particular kernel :
`cudaError_t cudaFuncSetCacheConfig(const void*
func, enum cudaFuncCacheConfig cacheConfig);`

3.2.1 Memory bound kernels

Register spilling

- When more register than max allowed per thread : spill to global memory (local memory on Keplers).

3.2.2. Instruction bound : toward occupancy sumMatrix example

- Goal : saturate instruction bandwidth and memory bandwidth
- Vocabulary : number of active wraps : parallelism exposed to an SM (streaming multiprocessor).
- Attention : high occupancy -> not always higher perf

sumMatrix.cu

- Tune the blocksize
 - Example in sumMatrix
 - `./sumMatrix <blockXsize> <blockYsize>`
- Try following combinations : [16 32], [32 32], [16 16] , [32 16]

Measuring occupancy

- `nvprof —query-metrics` to see a list of supported metrics
- Use `nvprof —metrics achieved_occupancy` to see the resulting occupancy for each cases

Load throughput

- `nvprof —metrics gld_throughput`
- Which combination has the highest throughput ?

Load efficiency

- `nvprof —metrics gld_efficiency`
- Which combination has the highest load efficiency ?

Other good practices

Exposing more parallelism

- Loop unrolling : gains up to 600%

« School case » example : reduction

http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf

The pdf is in SLIDES/DOCS folder
Also example at OPTIMIZATIONS/reduceInteger